



EVIDENCE ROOM

AP AI Readiness Diagnostic

Free lead magnet · 36 questions · six dimensions

Proof before permission.

Score the organisation before anyone commissions an unofficial copilot. No invoice upload. A5 at 0 or 1 vetoes placement above maturity L2.

Hard hold. ERP remains the system of record for accounting. Named humans remain the system of record for cash. Vendor-bank change and close attestation stay human. Duplicate & Anomaly flags are hypotheses, not fraud verdicts. This book is a method, not hosted software.

How to use this book

Buy Friday. Begin Monday. Do not charter from memory.

1. Score the free diagnostic (36 questions). Honour vetoes.
2. Map one path (usually PO-goods). Code exceptions from the closed list.
3. Charter Wave 1 at L0 or L1: Intake, Validation, Matching, Triage, Duplicate & Anomaly, Orchestrator.
4. Write the payment sentence on the wall. Agent 12 annotates; humans release.
5. Shadow before any write path. Promote only with an evidence pack.

Every page should answer at least one of: What should I do? How? Who owns it? What can go wrong? How do I control it? How do I measure it? What evidence proves it works?

Northline Industrials is a *fictional* manufacturer used so tables can be concrete (~15,000 invoices/month, Dynamics 365, Cleveland SSC). Figures are illustrative, not a client result.

Customer-facing statistics are limited to the research ledger. Do not paste vendor cost-per-invoice figures into a business case as targets.

AP Agent Readiness Diagnostic

SKU: ER-DIAG · \$0
Lockup: EVIDENCE ROOM — AP AGENT OS
Brand idea: Proof before permission.
Time: About 45–60 minutes with last-quarter facts.
Output: A maturity band (L1–L5), a heatmap, and a suggested first wave at observe/recommend.

This pack is a reading aid. It is not a purchase of the AP Agent OS. It does not certify readiness, compliance, investment return, fraud coverage, or accounting quality. It does not connect to your ERP. Do not paste production invoices, bank details, or employee records into any Evidence Room form.

What you do

1. Read `DIAGNOSTIC_GUIDE.md`. Answer all 36 questions. Score 0–4 with evidence you can reopen.
2. Place the organisation on `MATURITY_MODEL.md`.
3. Use `10_AGENT_OPPORTUNITY_OVERVIEW.md` to see which roles are design options — not a mandate to stand sixteen agents.
4. If you need language for a sponsor, use `BUSINESS_CASE_STARTER.md`. It will not write a savings guarantee.

What the score means

A high score on technology and a 0 on payment control is not “mostly ready.” It is a control gap with software nearby.

The suggested Wave 1 is always the same six names unless you lack owners: Invoice Intake, Invoice Validation, Matching, Exception Triage, Duplicate & Anomaly, Orchestrator — at L0 or L1.

What this pack refuses

- A single “AI readiness” number that hides a payment-hold failure.
- A recommendation to auto-release payments.
- A comparison that treats Northline Industrials (fictional) as your benchmark.
- A certificate you can show an auditor as sufficiency.

After you finish

Band	Sensible next licence
L1–L2	Starter (\$79) to name Wave 1 on paper this week
L2–L3	Professional (\$199) if one person will hold the full OS
L3+ and several pens	Team (\$499)
Unusual stack, multi-ERP, or political load	Custom Blueprint application

Site: evidenceroom.ai/diagnostic

Document control

Field	Value
Toolkit	Evidence Room — AP Agent OS
Object	Free diagnostic pack README
Status	Edition 1.0.0

AP Agent Readiness Diagnostic — guide

Product: Evidence Room — AP Agent OS

Audience: Finance Transformation lead, AP Manager, Controller, Internal Audit observer

Use: Score the *agent layer* you have — or the absence of one — across six dimensions.

Not: A certificate, an ROI model, a fraud assessment, or a compliance audit.

Northline Industrials appears as a fictional worked example. Do not copy its scores into a board pack as comparables.

1. How to sit the diagnostic

1. Appoint one scorer and one challenger. The challenger asks “show me the artefact.”
2. Answer every question. “Not applicable” is allowed only for Q-T4 (RPA/iPaaS) and Q-E6 (early-pay programme) if the object truly does not exist — score those as 0 and note “absent,” not “fine.”
3. Score **0–4** using the rubric below. Use the lowest score that is fully true. Partial credit is not a 3.
4. Record the artefact (report name, SOP ID, ticket, or “none”).
5. Average each dimension (six questions). Do **not** average the six dimensions into one corporate grade.
6. Apply the veto rules in §4 before you speak a band.
7. Place the organisation on MATURITY_MODEL.md.
8. Fill the heatmap and the first-wave recommendation in §7–§8.

Timebox: 60 minutes. If a question takes more than four minutes, score 0 or 1 and move on. Mystery is a score.

2. Scoring rubric (0–4)

Score	Meaning	Evidence test
0	Absent	No artefact. Recollection only. “We would know.”
1	Informal	A habit in one person’s head or a Slack thread. Cannot be reopened by a deputy.
2	Documented	A SOP, policy, or deck exists. Not clearly used on the last 20 invoices.
3	Used	Artefact exists and last-month samples show it in operation. Gaps are named.
4	Gated / sampled	Used, sampled, owned, and change-controlled. A miss has a written response.

Write the score and a five-word evidence note. Example: 2 – SOP-AP-14, unused since March.

3. The 36 questions

Dimension A — Process (how work moves)

A1. Coded exceptions. Are parked or rejected invoices coded from a closed list (not free-text only)?

Look for: exception taxonomy or equivalent; last 50 parks.

Northline illustration (fictional): 3 — codes exist; 18% still “other.”

A2. Named path. Is there a written invoice-to-pay path that distinguishes PO 3-way, PO 2-way, non-PO, recurring, and intercompany?

Look for: process map with object types.

Northline: 3 — four paths; intercompany informal.

A3. Owner on the clock. Does every open exception have a named human owner and an aging start date?

Look for: queue that can be sliced by person, not only by “AP.”

Northline: 2 — queue by team; aging by invoice date, not code date.

A4. Match policy is written. Can a new specialist apply 2-way / 3-way / GR-IR and tolerances without asking a neighbour?

Look for: tolerance table, not a hallway rule.

Northline: 3 — table in Dynamics; services ≤ \$10,000 are 2-way.

A5. Payment run is a procedure. Is proposal build, review, authorise, and transmit written as separate steps with named roles?

Look for: last run’s audit trail of two human approvers.

Northline: 4 — Tuesday/Thursday; Payments Lead + Treasury.

A6. Close is a checklist. Does period-end AP have a checklist with accrual candidates, cut-off, and attestation — not a heroic weekend?

Look for: last close pack.

Northline: 3 — checklist; accrual list still spreadsheet-built.

Process dimension score: average A1–A6. Veto: if A5 is 0 or 1, maturity cannot be placed above L2 until the payment procedure is written.

Dimension B — Data (whether a packet can exist)

B1. Invoice source completeness. Can you state, for the last month, what share arrived as PDF email, XML/EDI, portal, or paper — and which source lacks a retrievable image or file?

Look for: intake report.

Northline: 3 — portal + email measured; 4% paper at Monterrey unmeasured.

B2. Extract quality is sampled. Is field-level extract accuracy sampled against a gold label (or dual key), not assumed from vendor slides?

Look for: last sample sheet.

Northline: 1 — vendor dashboard only.

B3. Vendor master fitness. Can AP see a unique vendor, payee status, and bank-change date without a tribal expert?

Look for: master-data report; recent bank-change log.

Northline: 3 — dual control on bank; duplicates persist.

B4. PO and receipt availability. For PO invoices, are PO and GR/IR (or service entry) queryable in the same window the matcher uses?

Look for: a matcher working a live item in front of you.

Northline: 3 — D365; delayed GR from Birmingham warehouse.

B5. Work-object identity. Does an invoice, exception, statement line, or proposal line have a stable ID that survives a tool change?

Look for: ID in ERP plus any ticket tool.

Northline: 2 — ERP doc number; tickets renumber.

B6. Evidence reopen. Can a Controller or auditor reopen the source, extract, match worksheet, and approval for a named invoice from last month in one sitting?

Look for: do it once during the diagnostic.

Northline: 2 — source and approval yes; match worksheet no.

Data dimension score: average B1–B6. A 4 on B1 with a 0 on B6 means you capture and cannot inspect.

Dimension C — Controls (who may do what)

C1. Payment hold is real. Is payment authorisation human in **both** policy and the system that transmits the file?

Look for: last bank file; who clicked.

Northline: 4 — two humans; Positive Pay.

C2. Vendor bank dual control. Does a bank-detail change require two humans and leave a before/after record?

Look for: last five bank changes.

Northline: 4.

C3. DOA is applied, not edited, by helpers. If any bot or copilot routes approvals, can it edit the DOA table?

Look for: access listing. Score 0 if unknown.

Northline: 3 — workflow applies; IT can edit with change ticket.

C4. Duplicate / anomaly clearance. When a possible duplicate is raised, is clearance a named human outcome (clear / confirm / escalate / defer) rather than a silent ignore?

Look for: last 20 flags or the admission that flags have no register.

Northline: 1 — email alerts; no outcomes.

C5. Segregation still holds. Are capture/match, vendor master, and payment release still different human roles when a model is assisting?

Look for: SoD listing plus one week of user activity.

Northline: 3 — roles hold; temporary access at close is messy.

C6. Autonomy register. Is every live model, copilot, or bot listed with owner, object type, and level — including “we only use it to draft emails”?

Look for: the list. Absence is 0.

Northline: 0 — two informal copilots, unlisted.

Controls dimension score: average C1–C6. **Veto:** if C1 or C2 is 0, do not commission any agent above L0. If C1 is 0, do not open Agent 12.

Dimension D — Technology (the stack you will keep)

D1. System of record is named. Is it written that ERP is the accounting system of record, and that no chat window is?

Look for: architecture one-pager or the absence of one.

Northline: 2 — assumed; not written.

D2. Capture stack is known. Can you name the capture / IDR / portal products and whether they are in support?

Look for: vendor list.

Northline: 3.

D3. Workflow and match tools are known. Same test for match, park, and approval workflow.

Northline: 3 — native D365.

D4. Integration fabric. Is there a named way work moves (iPaaS, RPA, API, human re-key) and an owner for it?

Look for: one diagram. Score 0 if “the bot team knows.”

Northline: 2 — two RPA jobs; owner on leave.

D5. Access for any future agent is least-privilege in principle. Do service accounts exist as a pattern, or would a pilot share a specialist’s login?

Look for: IAM pattern. Shared login is 0.

Northline: 1 — shared mailbox; shared UI login on one RPA.

D6. Prompt / model change control. If anyone is already using a model on invoices, is a prompt change treated as configuration change?

Look for: ticket. If no model is in use, score 1 (informal readiness), not 4.

Northline: 0 — prompts in personal chats.

Technology dimension score: average D1–D6. High technology does not raise the maturity band if Controls veto.

Dimension E — Talent (names, not “the SSC”)

E1. AP Process Owner exists. Is there a named person who can accept or reject an Orchestrator charter?

Look for: role on an org chart *and* the person confirms they know.

Northline: 3 — AP Manager named; Orchestrator idea is new.

E2. Controls Lead exists. Named person for Duplicate & Anomaly clearance above threshold.

Northline: 2 — role exists; not briefed on flags-as-hypotheses.

E3. Match skill is not a single point of failure. If the best matcher is out for two weeks, does the path still run?

Look for: backup and a written match tree.

Northline: 2 — two skilled matchers; tree partly oral.

E4. Time to sit evidence. Have you reserved hours for sampling and charter work, or is this “on top of close”?

Look for: a calendar block in the next 30 days. None is 0.

Northline: 1 — transformation asked; no hours freed.

E5. Audit is invited as observer. Has Internal Audit been told they will sample packets, not operate agents?

Look for: a mail or a standing agenda slot.

Northline: 1 — not yet asked.

E6. Language discipline. Have specialists been told not to write “fraud,” “compliant,” or “the AI approved it” on tickets?

Look for: a note or huddle.

Northline: 0.

Talent dimension score: average E1–E6. If E1 or E2 is 0, Wave 1 owners are incomplete — charter only after names exist.

Dimension F — Economics (visibility, not a promise)

Score **visibility and honesty** of cost. A 4 does not mean the organisation will save money.

F1. Volume is known. Invoices per month, in-scope vs. out-of-scope, last two quarters.

Look for: a report, not a round number remembered from a town hall.

Northline: 4 — 15,000/month (illustrative).

F2. Fully loaded AP cost is known. Team cost the organisation already uses internally (do not invent a benchmark).

Northline: 3 — SSC cost known; plant liaison time not included.

F3. Exception rate is known. Share of invoices that do not post first pass, with a definition.

Northline: 3 — 28% (illustrative observational baseline).

F4. Cost per invoice is estimated in a range. Even a wide range with a method beats a vendor slide.

Northline: 2 — spreadsheet once; not refreshed.

F5. Layer cost is budgeted as cost. Tools, model, implementation, and exception minutes have a place to be recorded per thousand invoices.

Northline: 0 — no ledger.

F6. Success is not only a savings slide. Is there written permission to measure operating discipline (coverage, coded fails, sampled worksheets, dual-human release) even if cycle time does not move in 90 days?

Look for: a sponsor sentence. “We need 3x ROI to start” is score 0.

Northline: 1 — sponsor asked for a savings range; Transformation refused to invent one.

Economics dimension score: average F1–F6. This dimension never authorises a guaranteed-savings statement in the output pack.

4. Veto and cap rules

Apply after scoring, before the band.

Condition	Effect
C1 (payment hold) is 0	Cap maturity at L1. No Agent 12. L0 only.
C1 is 1 or C2 is 0	Cap at L2. Fix holds before L1 recommendations on payment-adjacent objects.
E1 or E2 is 0	Wave 1 incomplete. Name owners before commissioning.
A5 is 0–1	Cap at L2. Write the payment procedure.
F6 is 0 and sponsor demands a guaranteed ROI	Do not use this diagnostic to authorise spend. The product will not write the guarantee.
Any question scored from memory with artefact “none”	Treat as 0 in a steering pack, even if the scorer was optimistic.

Do not average away a veto.

5. Opportunity heatmap

Plot dimension averages. Use the labels, not a traffic-light that implies “green = automate.”

Dimension	Average (0–4)	Heat	Reading
Process		Cold / Warm / Hot	Cold 0–1.4: path is oral. Warm 1.5–2.9: maps exist. Hot 3.0–4.0: used and sampled.
Data			Cold: packets cannot be reopened. Hot: source-to-approval in one sitting.
Controls			Cold: do not promote. Hot: holds are real; register exists.
Technology			Cold: stack unknown. Hot: named, supported, least-privilege pattern.
Talent			Cold: no owners. Hot: owners, hours, audit observer, language.
Economics			Cold: volume unknown. Hot: costs visible; success is not only ROI theatre.

Heat is not priority to automate. A Hot technology / Cold controls cell means: stop adding models.

Northline illustration (fictional, not a target):

Dimension	Avg	Heat
Process	3.0	Hot
Data	2.3	Warm
Controls	2.5	Warm (C6 = 0 pulls it down)
Technology	1.8	Warm
Talent	1.5	Warm / near Cold
Economics	2.2	Warm

Reading: *Northline can write Wave 1 charters. It should not promote anything. It should list the informal copilots tomorrow.*

6. Maturity placement

Take the **lowest** of: (a) the simple average of the six dimension scores, (b) the cap from §4, (c) the qualitative fit in MATURITY_MODEL.md.

If you land here	Say this
L1 Undesigned	We do not have an agent layer. We may have tools.
L2 Informal activity	We have prompts or bots without charters or a register.
L3 Named and held	Wave 1 is named; payment is written as human; autonomy is L0/L1.
L4 Evidence-gated	Promotion uses gates; packets reopen; flags have outcomes.
L5 Recertified layer	Demotion works; cost ledger runs; L3 is rare and sampled.

Most mid-market organisations that have “started AI in AP” sit at L2. That is a usable finding.

7. Ten-agent opportunity (overview)

See 10_AGENT_OPPORTUNITY_OVERVIEW.md for the working page. In the diagnostic close, tick only what the heatmap supports.

Agent	Commission now?	Rule of thumb
01 Intake	If B1 $\geq$ 2	L0 if B2 $\leq$ 1
02 Validation	If A1 $\geq$ 2 or you will write codes this month	L1 only with a checklist
03 Matching	If A4 $\geq$ 2 and B4 $\geq$ 2	L1; no L3 talk
04 Triage	If A1 $\geq$ 2 and A3 $\geq$ 1	Codes first
10 Duplicate & Anomaly	Always at L0 if C4 $\leq$ 2	Never as a fraud engine
16 Orchestrator	If E1 $\geq$ 2	Work objects
05 Goods Receipt	Wave 2 if missing GR is a top code	Packets, not dummy GRs
07 Approval	Wave 2 if non-PO is material	Applies DOA; does not approve
08 Supplier Resolution	Wave 2 if supplier latency is material	Draft only
12 Payment Proposal Review	Only if C1 $\geq$ 3	Annotate only

Agents 06, 09, 11, 13, 14, 15 wait unless Professional is in use and Wave 1 is stable.

8. Business-case starter (language only)

Do not attach a savings range to the diagnostic result.

Permitted close paragraph (edit names):

The diagnostic places us at **L[n]**. Payment authorisation is **[human / not evidenced]**. I recommend chartering six agents at observe/recommend, naming owners, and running a two-week shadow on [path]. Success for the first 30 days is packet completeness, coded exceptions, recorded flag outcomes, and an autonomy register — not a cost-takeout number. Evidence Room — AP Agent OS is a method for that work. It does not replace our stack and does not promise savings, fraud detection, compliance, or ROI.

Full language: BUSINESS_CASE_STARTER.md.

9. Downloadable scorecard — instructions

Produce a local worksheet (spreadsheet or paper). Evidence Room does not host your scores.

Sheet tabs (recommended)

1. **Cover** — organisation, scorer, challenger, date, edition 1.0.0, “not a certificate.”
2. **Scores** — 36 rows: ID, question short name, score 0–4, artefact, note.
3. **Dimensions** — six averages, veto flags, cap.
4. **Heatmap** — table in §5.
5. **Band** — L1–L5 one sentence.
6. **Wave 1** — six agents, owners, backups, starting level.
7. **Holds** — C1–C2 evidence.
8. **Open questions** — anything scored from memory.

Column definitions for Scores

Column	Rule
ID	A1–F6
Score	Integer 0–4
Artefact	File name, SOP ID, report, or none
Recency	Date the artefact was last used
Challenger agree	Y/N
If N, lower score wins	Do not split the difference

What not to put on the scorecard

- A vendor's cost-per-invoice cited as yours.
- A fraud-loss estimate.
- A compliance score.
- Northline's numbers without the word *illustrative*.
- An ROI.

Print the cover and the band page for the Monday note. Leave the 36-row sheet in the working file.

10. After the diagnostic

Next artefact	Where
Maturity definitions	MATURITY_MODEL.md
Agent opportunities	10_AGENT_OPPORTUNITY_OVERVIEW.md
Sponsor language	BUSINESS_CASE_STARTER.md
Twelve steps	00_READ_ME/QUICK_START.md
Starter operating model	02_AP_AGENT_STARTER/

Document control

Field	Value
Toolkit	Evidence Room — AP Agent OS
Object	Diagnostic guide (36 questions)
Scoring	0–4; six dimensions; veto caps
Status	Edition 1.0.0
Not	Certification, audit, or financial advice

AP agent-layer maturity model (L1–L5)

Product: Evidence Room — AP Agent OS

Use: Place the organisation after the 36-question diagnostic.

Rule: The band describes the *agent layer*, not the quality of the ERP and not the talent of the AP team.

A strong shared-service team on a modern ERP can sit at L1 on this model. That is not an insult. It means agents are undesigned. The stack can still be excellent.

This model does not measure fraud resilience, regulatory compliance, accounting accuracy, or investment return.

1. The five levels

Level	Name	One sentence
L1	Undesigned	Tools and people exist. No agent layer exists.
L2	Informal activity	Models or bots are in use without charters, a register, or evidence gates.
L3	Named and held	Wave 1 is chartered at L0/L1. Payment and vendor-bank holds are written and real.
L4	Evidence-gated	Promotion uses gates. Packets reopen. Flags have recorded human outcomes.
L5	Recertified layer	Demotion and retirement work. Cost of the layer is visible. L3 is rare, sampled, and still not payment.

Autonomy levels inside a charter (Observe / Recommend / Prepare / Execute / Managed) are a different scale. Do not say “we are L4” when you mean an agent was promoted to execute-within-guardrails. In conversation, say **maturity L4** or **autonomy L3**.

2. Definitions that can be tested

L1 — Undesigned

True when most of these hold:

- No named AP agents. “We might use Copilot” is not a name.
- Exceptions are free-text or oral.
- No autonomy register.
- No work-object identity beyond the ERP document number and a shared mailbox.
- Payment may still be well controlled. That control is *AP hygiene*, not agent-layer maturity.

Monday move: Sit the diagnostic. Do not buy a model licence to “get to L2.” L2 is not progress if it is only informal activity.

Northline (fictional) analogue: Northline in 2024 before anyone opened a chat on an invoice.

L2 — Informal activity

True when:

- At least one model, copilot, RPA job, or vendor “AI feature” touches invoices, exceptions, or supplier mail.

- There is no charter, or the charter is a slide.
- Prompts live in personal accounts.
- Flags, if any, can be ignored without a recorded outcome.
- Someone has said “it just helps” in a steering meeting.

Risk: Informal activity is how payment-adjacent features get turned on in a demo and left on.

Monday move: List every informal helper. Charter Wave 1 at autonomy L0. Turn *off* anything that can post, send externally, or annotate a payment proposal until a charter exists.

Northline (fictional): Capture copilot + matching suggestions; unlisted; C6 = 0.

L3 — Named and held

True when:

- Six Wave 1 agents have owners, object types, exclusions, and a starting autonomy of Observe or Recommend.
- Payment authorisation and vendor-bank change are human in policy **and** system.
- Orchestrator work objects exist on the pilot path.
- Internal Audit has a date to observe packets.
- No agent is commissioned at autonomy L3 or L4.

Monday move: Run the two-week shadow. Refuse scope expansion. Write the cost ledger, even if the numbers are small.

L4 — Evidence-gated

True when L3 is true and:

- Promotion cases use written gates (coverage, completeness, sampled disagreement, no silent scope, failure handling, cost meter, language, owner).
- A Controller or auditor can reopen source, extract, worksheet or coded fail, and approval for a sampled invoice in one sitting.
- Agent 10 outcomes are recorded (clear / confirm / escalate / defer).
- Prompt and limit-table changes are change-controlled.
- At least one demotion or declined promotion has occurred — or a tabletop of demotion has been tested. A culture that cannot demote is not gated; it is optimistic.

Monday move: Open Wave 2 prepare-level packets only for named object types. Agent 12 remains annotate-only.

L5 — Recertified layer

True when L4 is true and:

- Every live agent has a recertification date. Overdue means automatic demotion.
- Cost per thousand invoices for the layer (licence + model + exception minutes) is reported with the operational pack.
- L3 execute-within-guardrails exists only on a named class with a limit table and sampling.
- Payment release is still dual-human. This remains true at L5. There is no L6 that pays.

- Root Cause proposals are accepted or rejected by process owners; they are not silently implemented by a model.

Monday move: Recertify. Publish the register. Do not celebrate “touchless.”

3. Placement method

1. Score the diagnostic.
2. Apply veto caps (DIAGNOSTIC_GUIDE.md §4).
3. Read the qualitative tests above. If the tests for L4 fail, you are not L4 even if the arithmetic average is 3.2.
4. When in doubt, take the lower band.

Arithmetic average of dimensions	First guess	Override
0.0–1.4	L1	Informal bots in use → L2
1.5–2.4	L2	Wave 1 chartered and holds real → L3
2.5–3.2	L3	Gates and reopen test fail → stay L3
3.3–3.7	L4	Cannot demote / no flag outcomes → L3
3.8–4.0	L5	Any payment-adjacent execute → not L5; control incident

4. What does not move the band

- Buying Professional or Team. A licence is not maturity.
- Invoice volume. Small books can be L4; large books can be L1.
- ERP brand.
- A vendor’s “AI included” SKU.
- A single impressive demo on a clean PDF.
- A promised savings model. Economics visibility can score high while the layer is still L2.

5. Language for the sponsor

Use:

We are at maturity L2 (informal activity). That is a control finding about undocumented helpers, not a finding that AP is weak. The next band is L3: named Wave 1, human payment hold, autonomy register.

Do not use:

We are only L2, so we are losing millions / we are non-compliant / we are exposed to undetected fraud.

Those sentences reintroduce banned claims.

6. Re-sit cadence

Trigger	Action
90 days after Wave 1 shadow starts	Re-sit dimensions C, E, and the reopen test (B6)
Any payment-adjacent incident	Immediate re-sit of C; cap band until closed
New ERP company code or new capture vendor	Re-sit B and D for that path
Request to promote an agent	Band must be L3+ and gates must pass; the band does not replace the gate pack

Keep the prior scorecard. Do not overwrite history.

Document control

Field	Value
Toolkit	Evidence Room — AP Agent OS
Object	Maturity model L1–L5
Related	DIAGNOSTIC_GUIDE.md
Status	Edition 1.0.0

Ten-agent opportunity overview

Product: Evidence Room — AP Agent OS
Audience: Diagnostic taker; Transformation lead choosing a first wave
Use: See ten roles as *design options*. Tick only what owners and holds support.
Not: A mandate to stand ten agents, a software bill of materials, or a savings menu.

The full operating system specifies sixteen agents. This free overview covers the ten that appear most often in a first two waves. Full specifications ship in Professional (03_AP_AGENT_OS_PRO/Agent_Library/). Starter contains working blueprints for the same ten (02_AP_AGENT_STARTER/TOP_10_AGENT_BLUEPRINTS.md). Northline Industrials is fictional. Its “fit” notes are illustrations.

How to read a row

Field	Meaning
Default autonomy	Where a new charter starts. Never L3.
Human owner	A role, then a name. “The SSC” is not an owner.
Opportunity	The <i>work object</i> that becomes inspectable — not a promised dollar.
Do not	Hard exclusion. If a vendor demo contradicts it, the demo loses.
Fit test	Diagnostic signals that the option is even discussable.

1. Invoice Intake (Agent 01)

Job	Capture, classify, extract, register inbound invoices.
Default autonomy	L1, or L0 if extract quality is unsampled.
Typical owner	AP Operations Lead
Opportunity	Every inbound invoice becomes a work object with a source and an extract.
Do not	Decide match. Post. Create vendors.
Fit test	You can name intake sources ($B1 \geq 2$).
Northline (fictional)	Email + portal; paper at Monterrey is a separate object type, not silently in-scope.

2. Invoice Validation (Agent 02)

Job	Test completeness, supplier identity, tax <i>fields</i> , posting readiness.
Default autonomy	L1
Typical owner	AP Quality Lead
Opportunity	Fail reasons become codes. The mix becomes analysable.
Do not	Decide tax position on ambiguous or cross-border items. Post.
Fit test	You will maintain a checklist (A1 or a decision to write codes this month).
Northline	US sales/use, GST/HST, IVA — flags only; Tax desk decides.

3. Matching (Agent 03)

Job	Apply 2-way / 3-way / GR-IR logic and the written tolerance table.
Default autonomy	L1
Typical owner	AP Match Lead
Opportunity	Inspectable worksheets on a sample; clean vs. break is explicit.
Do not	Invent a GR or a PO. Edit tolerances. Force-match.
Fit test	Match policy written ($A4 \geq 2$) and PO/GR queryable ($B4 \geq 2$).
Northline	3-way inventory; 2-way contracted services $\leq$ \$10,000 (illustrative policy).

4. Exception Triage (Agent 04)

Job	Classify the break, assign owner, set next action and SLA.
Default autonomy	L1
Typical owner	Exception Desk Lead
Opportunity	Parks stop being a pile. Aging is by owner and code.
Do not	Resolve the break. Waive PO. Approve price.
Fit test	A code list exists or will exist before go-live.
Northline	EX-MPO, EX-PRM, missing GR dominate the fictional mix.

5. Goods Receipt (Agent 05)

Job	Detect missing or aged receipts; assemble a receiver packet.
Default autonomy	L1; L2 prepare only when chartered.
Typical owner	Plant / warehouse finance liaison
Opportunity	Receivers see facts, not a nag. Dummy GRs become a visible refusal.
Do not	Create a GR to make an invoice match.
Fit test	Missing GR is a material code. Wave 2, not Wave 1.
Northline	Birmingham warehouse silence is the worked example.

6. Approval (Agent 07)

Job	Route non-PO and over-tolerance items per DOA; chase complete packets.
Default autonomy	L1
Typical owner	AP Approvals Coordinator
Opportunity	Packets wait for a deputy in the table, not a convenient manager.
Do not	Approve. Invent a deputy. Edit DOA.
Fit test	Non-PO share is material; DOA table exists.
Northline	Cost-centre owner on leave → wait; do not climb the tree.

7. Supplier Resolution (Agent 08)

Job	Draft supplier queries; track replies.
Default autonomy	L1; send remains human or human-approved.
Typical owner	Vendor Master / AP supplier desk
Opportunity	Queries contain facts. Tone is request, not concession.
Do not	Agree price, quantity, due date, or a pay date. Change bank details.
Fit test	Supplier latency is a measured delay, not a vibe.
Northline	“Please confirm the PO number” — never “we will update the PO.”

8. Duplicate & Anomaly (Agent 10)

Job	Flag possible duplicates and unusual patterns for human review.
Default autonomy	L0
Typical owner	AP Controls Lead
Opportunity	Flags become a register with four human outcomes.
Do not	Say “fraud.” Clear a high-value flag. Hold or release payment on its own.
Fit test	Always discussable. Commission at L0 even if C4 is weak — especially then.
Northline	0.8% suspect rate is illustrative observation, not a detection claim.

9. Payment Proposal Review (Agent 12)

Job	Annotate a payment proposal: holds, open flags, yesterday's bank changes.
Default autonomy	L1; ceiling L2 annotate.
Typical owner	Treasury / AP Payments Lead
Opportunity	Humans see a marked proposal. The release click stays human.
Do not	Approve, deselect-and-release, transmit, or "pay to keep the discount."
Fit test	$C1 \geq 3$ (payment hold real). Otherwise do not open this agent.
Northline	Tuesday run, 612 lines, two uncleared flags — humans deselect (illustrative).

10. Orchestrator (Agent 16)

Job	Sequence work, enforce SLAs, record handoffs and evidence URIs.
Default autonomy	L1
Typical owner	AP Process Owner
Opportunity	A work object for every invoice, exception, statement line, and proposal line on the path.
Do not	Post, pay, change vendors, or pick an approver outside DOA.
Fit test	$E1 \geq 2$ (process owner exists).
Northline	System of record for <i>state</i> ; D365 remains system of record for <i>accounting</i> .

Six you are not being asked to stand from this pack

ID	Agent	Why it waits
06	PO Quality	Procurement-owned; Wave 3 diagnostic
09	Internal Follow-up	Needs triage codes and packets first
11	Vendor Statement	Periodic; easy to confuse with "the statement is the invoice"
13	AP Close	Attestation stays human; easy to over-scope
14	AP Reporting	Only as good as definitions
15	Root Cause	Needs a coded history

Professional contains their specifications. Curiosity is not a Wave 1 plan.

Choosing among the ten

Use this order, not a heat map of "value."

1. If payment hold (C1) is weak → only 01–04, 10, 16 at L0. No Agent 12.
2. If owners E1/E2 are missing → write names before any charter.
3. If match policy or PO/GR data is weak → Intake + Validation + Orchestrator only, until A4/B4 move.
4. Otherwise → standard Wave 1 (01, 02, 03, 04, 10, 16).
5. Add 05, 07, 08 only after two weeks of shadow and a coded mix.

6. Add 12 only when C1 is used, not merely documented.

Document control

Field	Value
Toolkit	Evidence Room — AP Agent OS
Object	Ten-agent opportunity overview (free)
Related	Diagnostic; Starter blueprints; Professional library
Status	Edition 1.0.0

Business-case starter language

Product: Evidence Room — AP Agent OS

Audience: Transformation lead, AP Manager, Controller writing a short sponsor note

Use: Paste and edit. Replace brackets. Delete any sentence you cannot defend from your own diagnostic.

Not: A savings model, an ROI calculator, or Evidence Room's forecast of your results.

Professional contains a full model you can populate with *your* inputs (03_AP_AGENT_OS_PRO/Business_Case/). This file is language only.

1. Rules for this note

1. Lead with the layer you lack, not the money you hope to free.
 2. Name the human holds in the same paragraph you mention agents.
 3. If you cite external research, use the labelled citations in §6 and no others.
 4. Do not attach a payback period.
 5. Do not use Northline's figures as yours. Northline Industrials is fictional.
 6. If the sponsor requires a guaranteed savings number as a condition of starting, this product is the wrong artefact to put in the pack.
-

2. Subject line options

- AP agent layer — diagnostic reading and proposed Wave 1
- Observe/recommend agents in AP — named owners, human payment hold
- Request: two-week shadow on [path], not a platform replacement

Avoid: "AI ROI," "touchless AP business case," "fraud reduction case."

3. One-page note (canonical)

To: [CFO / CAO / VP Finance]

From: [name, role]

Date: [date]

Re: Accounts Payable agent layer — readiness reading

We were asked to consider AI in Accounts Payable. I sat a structured diagnostic of the *agent layer* — whether we have named roles, inspectable evidence, a permission ladder, and a written human hold on payment — not a review of our ERP.

Reading. We are at **maturity [L1–L5 — name]**. [One sentence from the diagnostic band.] This is not a finding that the AP team is weak and not a compliance certificate.

Holds. Payment authorisation and vendor bank-detail change [are / are not] evidenced as human in both policy and system. I will not recommend any helper that releases a payment file. Duplicate and anomaly signals, if we raise them, will be hypotheses for [Controls Lead name], not fraud findings.

Proposal. Charter six agents at observe or recommend on [path / entity]: Invoice Intake, Invoice Validation, Matching, Exception Triage, Duplicate & Anomaly, Orchestrator. Owners: [six names]. For two weeks the agents write shadow records only. Operators work as they do today.

Success at 30 days. Every in-scope invoice has a work object; validation fails are coded; a sample of match worksheets can be reopened; flags have a recorded human outcome; the autonomy register lists every live helper at observe/recommend; payment release still follows [our existing dual-approval procedure]. I am not proposing a cost-takeout target.

What we would licence. [Starter \$79 / Professional \$199 / Team \$499] — a method and toolkit (Evidence Room — AP Agent OS). It does not replace [ERP / capture / payment tool]. It does not connect to our systems. Implementation stays with us.

What I am not asking. Permission to auto-post, auto-pay, or report a guaranteed return.

Ask. [Approve the two-week shadow / name the Controls Lead / book Internal Audit as observer on day 30 / approve the licence].

4. Extra paragraphs you may add — and those you may not

Allowed

Stack stays. We keep [ERP]. We keep [capture vendor]. We keep [payment procedure]. The work is to design the layer that sits across them.

Informal risk. We already have [copilot / RPA / vendor AI feature] touching invoices without a charter. The diagnostic treats that as informal activity. Listing it is the first control action.

Industry context (optional, labelled). Independent AP research is sometimes cited by vendors: Ardent Partners *State of ePayables 2024*, as cited by Tipalti, reported best-in-class electronic cost per invoice of \$2.78 versus \$12.88 for laggards, and cycle times of 3.1 versus 17.4 days. Ardent *AP Metrics that Matter 2025*, as cited by Tipalti, reported an exception-rate comparison of 22% to 9%. These are **vendor-originated citations of independent research**. They are not our baseline and not a target I am committing to.

Adoption context (optional, labelled). McKinsey's November 2025 finance AI work reported that 44% of 102 CFOs used gen AI for five or more use cases (from 7%), and that 65% were increasing gen AI investment. McKinsey *State of AI 2025* (survey June–July 2025, n=1,993) reported 88% regular AI use, roughly two-thirds not scaling, 62% experimenting with agents, 23% scaling an agent somewhere in the enterprise, and 39% reporting any EBIT impact. Workflow redesign is noted as a high-performer differentiator. I cite this as context for *why the question landed on my desk*, not as evidence that an AP agent layer will move EBIT.

Not allowed

- “This will save \$X.”
 - “Payback in Y months.”
 - “We will detect fraud.”
 - “This makes us SOX compliant / audit-ready.”
 - “The agent will be accurate.”
 - “We will move to autonomous payments.”
 - “Northline achieved ...”
-

5. Email to Internal Audit (observer, not operator)

I am proposing a two-week shadow of six named AP agents at observe/recommend. Payment release remains [our procedure]. I would like Audit to observe packets at day 30 — source, extract, worksheet or coded fail, flag outcome — not to operate the agents. Flags will be labelled as hypotheses. I am not asking you to certify the design.

6. Citation block (copy exactly if used)

Sources (not our results). Ardent Partners, *State of ePayables 2024*, as cited by Tipalti: \$2.78 vs \$12.88 cost per invoice; 3.1 vs 17.4 days. Ardent Partners, *AP Metrics that Matter 2025*, as cited by Tipalti: exception rate 22% → 9% in that comparison. McKinsey, November 2025, finance AI: 44% of 102 CFOs, five or more gen AI use cases (from 7%); 65% increasing investment. McKinsey, *State of AI 2025*, survey June–July 2025, n=1,993: 88% regular use; ~2/3 not scaling; 62% experimenting with agents; 23% scaling an agent somewhere; 39% any EBIT impact; workflow redesign as a high-performer differentiator. All figures are third-party. None is a forecast for this organisation.

7. If they ask “what’s the ROI?”

Reply:

I will not put a return number on this request. We can, later, record cost per thousand invoices for the layer and, if we choose, build a model from our own baseline. Evidence Room’s professional pack includes a calculator specification that uses our inputs and three cases — conservative, base, upside — and still does not promise the output. If a number is required to start, it must be ours, labelled as a management judgement, and it cannot be a condition I ask Evidence Room to guarantee.

Then stop talking.

Document control

Field	Value
Toolkit	Evidence Room — AP Agent OS
Object	Business-case starter language (free)
Related	Diagnostic; Professional Business_Case/
Status	Edition 1.0.0

Disclaimer — Draft for Counsel

Status: Working draft. **Not legal advice.** Intended for /disclaimer, checkout footers, and the zip. Counsel must confirm jurisdiction-specific wording.

Last updated: 20 September 2026

Lockup: EVIDENCE ROOM — AP AGENT OS

If this Disclaimer conflicts with a marketing sentence, **this Disclaimer controls** and the marketing sentence should be removed.

1. Nature of the product

Evidence Room — AP Agent OS is a **digital method and toolkit**. It contains written architectures, charters, frameworks, and examples to help an organisation **design, govern, and discuss** an Accounts Payable agent layer.

It is **not**:

- software that connects to, posts to, or operates your ERP
 - a payment service, money-transmission service, or bank
 - an audit, assurance, or attestation engagement
 - legal, tax, accounting, or regulatory advice
 - a certification that you meet any law, standard, or control framework
 - a fraud-detection, AML, or sanctions-screening product
 - a guarantee of any financial or operational result
-

2. Outcomes we do not promise

Without limitation, we do **not** promise or guarantee:

- savings, cost reduction, or **ROI**
- **fraud detection**, fraud prevention, or the absence of duplicate or abusive payments
- **regulatory compliance** (including but not limited to SOX, VAT/GST rules, employment, or data-protection law)
- **accounting accuracy** or correct postings
- **autonomous payment safety**
- cycle-time reduction
- “touchless AP” as an operating state
- that any model you later use will be correct

Any number in the materials labelled illustrative, fictional, or Northline Industrials is **not** a benchmark and **not** a target.

3. Payment authority

The method states that **payment authorisation and release remain human**. That is a design rule in a document. It does not itself stop a person or a vendor from wiring a model to a payment file. **You** must implement holds in your systems and policies.

We are not liable for payments you or your agents release.

4. Implementation is yours

You choose vendors, models, prompts, permissions, owners, and whether to commission any agent. You are responsible for testing, sampling, access control, and local policy.

The toolkit cannot see your invoices unless **you** paste them somewhere. Do not paste production data to us.

5. No professional relationship

Using the site or the files does not create an attorney-client, auditor-client, fiduciary, or similar relationship. Blueprint facilitation, if any, is scoped method design, not an audit opinion.

6. Third parties and names

Mentions of ERP or AP vendors are descriptive. No sponsorship or compatibility certification is implied unless a written agreement says so (none at the time of this draft).

We are **not** affiliated with Evidence Room LLC (forensic animation / evidence-room.net). See Trademark notes.

Lemon Squeezy and Stripe are independent. Their MoR and payment terms are theirs.

7. Site content

Resources, charts, and posts are educational and commercial description. They are not advice to buy or refrain from buying any security or to take any tax position.

8. Limitation pointer

Liability, to the extent it can be limited, is addressed in the Terms. **Counsel must write that limit.** Nothing here is an additional warranty.

9. Short footer (site and checkout)

Evidence Room — AP Agent OS is a digital method and toolkit. It is not an ERP, payment service, audit, or compliance certification. It does not detect fraud, authorise payments, or guarantee savings, accuracy, or ROI. You remain responsible for implementation and for all payment decisions.

Counsel checklist

- ☐ Conform to mandatory consumer wording
- ☐ Harmonise with insurance if any
- ☐ Confirm “not advice” is not over-claimed in Blueprint sales calls