

TIER 0 · FREE LEAD MAGNET

AP AI Readiness Diagnostic

Proof before permission. Thirty-six questions. A score you can defend.

Evidence Room — AP AI Readiness Diagnostic

Product: Evidence Room AP Agent OS — Tier 0 (Free Diagnostic)

Brand: Evidence Room — AP Agent OS

Brand idea: Proof before permission.

Line: Agents earn responsibility. Evidence decides.

Audience: CFO, Controller, Head of AP, Internal Audit observer

Time to complete: 30–40 minutes

ERP stance: Agnostic. Map answers to your live system names.

Examples: ACME Manufacturing — **ILLUSTRATIVE** only. Not a client. Not a case study.

Version: 1.0

Classification: Self-administered diagnostic. Not a control attestation, not a savings estimate, not an audit opinion.

Cover and positioning

Most finance teams have already bought something that extracts invoices, routes approvals, or “uses AI.” That is not the question this diagnostic answers.

The question is: **can this AP operation produce evidence that would justify giving an agent any responsibility at all?**

Evidence Room treats agents as tools with job descriptions. They start at observe. They move only when a named human can show precision, miss rate, exception cost, and a kill-switch. This file is the first measurement of whether that discipline is even possible in your shop.

What this diagnostic is for

- A CFO who will trade an email address for a clear picture of AP readiness — not a vendor demo.
- A Controller who needs to know whether the control environment would survive an agent touching invoices.
- An AP Manager who must pick one bounded first agent, not sixteen.

What you will walk away with

1. A transparent score out of 144, mapped to a five-level maturity model.
2. Six domain scores so a high “AI enthusiasm” score cannot hide a weak control or data score.
3. An opportunity heatmap you plot from *your* volume, pain, control risk, and data readiness — not from a vendor ranking.
4. A first-wave view of ten agents (the subset of the sixteen-agent stack you would consider first).
5. A baseline KPI worksheet with blanks. You fill them. We do not invent industry targets.
6. A business-case starter that structures conservative / base / upside *inputs*. It does not compute or promise savings.

Voice of this product. Intelligent, concise, executive, specific. No hype. No guaranteed savings, fraud detection, regulatory compliance, accounting accuracy, autonomous payment safety, or ROI. If a number is not one of the three researched facts cited below, or a figure you enter from your own ledger, it is labelled **ILLUSTRATIVE**.

Context you may quote (researched; not Evidence Room performance)

These facts describe the market, not this product:

- **APQC (16 Mar 2026):** top performers spend about \$0.38 per \$1,000 revenue to process AP vs about \$0.92 for bottom performers. Source: [APQC — How organizations can reduce accounts payable costs](#).
- **Deloitte Finance Trends 2026 (8 Oct 2025):** 63% of finance teams have fully deployed and actively use AI; only 21% report clear, measurable ROI; 14% are using fully integrated AI agents. Source: [Deloitte Finance Trends 2026 survey release](#).
- **The Hackett Group (19 Nov 2025):** AP *solution customers* averaged 60% touchless invoice processing; organisations with ≥30% touchless had 3.5× higher AP productivity; cycle times improved 59% post-implementation. This is **software-adopter research**, not Evidence Room performance. Source: [Hackett — AP solutions delivering gains through AI](#).

Use them as context for a board conversation about *discipline*. Do not paste them into a business case as “our expected outcome.”

How to complete (30–40 minutes)

What to do. Score 36 questions across six domains. Enter your own baseline KPIs. Plot three opportunities on the heatmap. Stop. Do not design an agent in this sitting.

How.

1. Appoint a **scorer** (Head of AP or AP Manager) and a **challenge** (Controller or Control owner). One person scoring alone produces optimism.
2. Open this file and, if you have it, the companion workbook described in `SCORECARD_INSTRUCTIONS.md`.
3. Work domain by domain. For each question pick **one** integer 0–4 using the rubric. If you are between two scores, take the lower one.
4. Write a one-line evidence note: the report, queue, SOP version, or “we do not have this.” A score without a note is a guess. Guessing is allowed — mark it `G`.
5. Do not workshop the answers. If two people disagree by two or more points, record both and use the lower score. The disagreement *is* the finding.
6. Complete the baseline KPI worksheet with blanks or “unknown.” Unknown is a valid answer and lowers Data & systems and Agent readiness when those questions ask about measurement.
7. Plot the heatmap after scoring, not before. The score tells you whether you are allowed to get excited about an opportunity.
8. Timebox: 25 minutes on questions, 8 minutes on KPIs and heatmap, 5 minutes on interpretation. If you overrun, you are designing, not diagnosing.

Who.

Role	Duty
Head of AP (scorer)	Answers process, exception, talent, and agent questions
Controller or Control owner (challenge)	Answers control and audit questions; vetoes inflated scores
Finance Systems (optional, 10 minutes)	Confirms data, interface, and identity facts
CFO (recipient)	Reads the interpretation page, not the 36 rubrics

What can go wrong.

- Scoring the SOP instead of the floor. The rubric is about what operators *do*.
- Letting a capture-vendor dashboard substitute for AP's own evidence.
- Scoring "we bought AI" as Agent readiness. A licence is not a charter.
- Completing it after a vendor demo. Recency bias will push every score up one point.

Control. Challenge initials the cover. Any domain averaging ≥ 3.0 must cite artefacts, not memory. Guess-marked items cannot be the sole support for a "ready for agents" conclusion.

Measure. Minutes to complete; count of **G** (guess) notes; number of two-point disagreements.

Evidence. Dated score sheet (this file or the workbook), initials, and the baseline KPI page. Store under [BUYER]/Evidence/Diagnostic/.

ACME — ILLUSTRATIVE. Head of AP and Controller sit for 38 minutes. They disagree on CA-03 (duplicate control) by two points; they take 1. Twelve answers are marked **G**. The CFO receives a one-page interpretation, not a 16-agent roadmap.

Scoring method (transparent, not a black box)

What to do. Convert 36 integer answers into six domain scores and one overall score. Map the overall average to a maturity level. Do not invent weights.

How.

Object	Formula
Question score	Integer 0, 1, 2, 3, or 4. No halves.
Domain score	Sum of its 6 questions $\div$ 6. Report to one decimal.
Overall score	Sum of all 36 questions $\div$ 36. Report to one decimal.
Overall points	Sum of all 36 questions. Maximum 144.
Maturity level	Look up overall score in the maturity table below.

There are **no hidden weights**. Process integrity does not outrank Controls. A shop that scores 3.8 on Agent readiness and 1.1 on Controls is not "almost ready." It is Experimental, because the

overall average will say so — and because the interpretation rules treat a Controls domain below 2.0 as a **hard brake** on any execute-level ambition.

Confidence tag (optional, does not change the score).

Tag	Meaning
E	Evidence on file (report, extract, SOP version, access list)
0	Observed on the floor in the last 90 days
G	Guess / tribal knowledge

If more than 12 answers are **G**, the diagnostic is **indicative only**. Re-run after a two-hour evidence pull. Do not take an indicative result to a steering committee as a decision.

Who. Scorer computes. Challenge checks arithmetic and the hard-brake rule.

What can go wrong. Averaging only the high domains. Re-weighting “because AI is strategic.” Converting 4s into a percentage that looks like an exam grade.

Control. Publish all six domain scores next to the overall score. Never publish overall alone.

Measure. Domain range (max – min). A range ≥ 2.0 is a finding: the operation is uneven, not “average 2.4.”

Evidence. The arithmetic is this section. A workbook that changes weights is not this diagnostic.

The 36 questions

Score each item 0–4 using the AP-specific rubric. “Why it matters” is the reason a CFO should care — not a prompt to buy software.

Domain 1 — Process integrity

Does a second person know how an invoice actually moves — and would an agent be automating the real path or the poster?

PI-01 — Documented as-done path

Question. Is there a current, signed description of how invoices move from receipt to ready-for-payment *as operators do it*, including informal queues?

Score	What it looks like in AP
0	No SOP, or the SOP is a vendor implementation leftover. Operators cannot point to it.
1	SOP exists; last review > 18 months; floor uses email folders and side spreadsheets the SOP omits.
2	SOP current for the happy path. Exception and payment-prep paths are tribal.
3	As-done map covers every live channel and the top exception types. Informal workarounds are listed, not hidden.

Score	What it looks like in AP
4	As-done map versioned, signed by Head of AP, sampled by Controls in the last 12 months. Floor and document match on a live walkthrough.

Why it matters. An agent that follows the SOP while clerks follow Jane's spreadsheet will create a second, uncontrolled path.

PI-02 — Channel inventory

Question. Can you name every inbound channel (mailbox, portal, EDI/cXML, scan, punch-out, shared inbox) and the volume on each for the last complete quarter?

Score	What it looks like in AP
0	"Invoices come in by email." No mailbox list. No volume split.
1	Main mailbox known. Portal and EDI exist but are not counted. Plant inboxes are unofficial.
2	Channels listed. Volume is a guess or a vendor dashboard that does not reconcile to ERP receipts.
3	Channel list with quarterly counts that reconcile to invoices received in ERP / workflow ($\pm$ a documented residual).
4	Channel list owned, volume reconciled, and each channel has a named intake owner and a fallback if that channel dies.

Why it matters. First-wave agents fail on the channel nobody listed. Volume without a channel split makes a heatmap fiction.

PI-03 — PO vs non-PO discipline

Question. Is it unambiguous, in policy *and* in the queue, which invoices require a PO, and what happens when they do not have one?

Score	What it looks like in AP
0	"Most things should have a PO." Clerks decide case by case.
1	Policy exists. After-the-fact POs are routine and unflagged.
2	PO-required categories written. Non-PO path exists but is used as a bypass.
3	PO-required flag is system-enforced by category or vendor. Retro-POs are flagged and reported.
4	Policy, system flag, retro-PO report, and a named Procurement owner for repeat bypass. Non-PO is a designed path, not a leak.

Why it matters. Missing-PO is usually a high-volume exception. If policy is soft, an agent will either invent a PO or rubber-stamp a bypass.

PI-04 — Match rule is written

Question. Are 2-way / 3-way / tolerance rules written as a table an auditor could test — not "we match when it looks right"?

Score	What it looks like in AP
0	Matching is judgement. Tolerances live in someone's head.
1	ERP has a tolerance. Nobody can produce the current values and who last changed them.
2	Tolerance table exists. Service POs, blanket POs, and multi-line invoices are handled by exception folklore.
3	Documented match policy: which invoice types are 2-way vs 3-way, numeric tolerances, who may widen them, and that widening is a change — not a clerk action.
4	Policy version cited on every exception. Config export matches the document. Changes go through a change record.

Why it matters. A language model must not be the match engine. If the rule is unpublished, you cannot tell whether the agent or the human is guessing.

PI-05 — Exception language

Question. Do parked invoices use a finite, shared set of reason codes, or a comment field?

Score	What it looks like in AP
0	Free-text holds. "Waiting," "check," "supplier."
1	ERP reason codes exist. >30% of volume is Other / Misc / blank.
2	Codes used on most items. Same break has three names across teams.
3	Published taxonomy (or a mapped ERP list). Primary code required. Misc has a weekly cap review.
4	Taxonomy versioned; ERP mapped 1:1; reports run by code, owner, and age — not by "exceptions."

Why it matters. Agents route by code. A comment field cannot own an SLA.

PI-06 — Close and cutoff are designed

Question. Is period-end for AP a checklist with owners (accruals, GR/IR, unrecorded spend, held payments), or a heroic weekend?

Score	What it looks like in AP
0	Close is "get through the pile." No checklist.
1	Checklist exists. Half the items are informal. GR/IR is a surprise every month.
2	Checklist used. Accruals depend on one person. Statement gaps are optional.
3	Named close owner, checklist, GR/IR ageing reviewed, unrecorded-spend scan defined.
4	Close pack retained; cutoff rules written; agent-touched items cannot skip the checklist.

Why it matters. An agent that posts faster into a sloppy cutoff moves the error into the next period. Speed is not close quality.

Domain 1 evidence note. Artefacts that support a high score: as-done map, channel volume extract, PO-required policy, match-tolerance export, reason-code mix, close checklist.

Domain 2 — Data & systems

Can you read a trustworthy invoice, vendor, PO, GR, and payment record — and write evidence — without inventing fields?

DS-01 — Invoice identity

Question. Is there a stable key for “this invoice” across channel, workflow, and ERP (supplier + invoice number + entity, or equivalent), and do you know how duplicates of the *key* are treated?

Score	What it looks like in AP
0	Invoice number is whatever was typed. Spaces and prefixes create new documents.
1	ERP has a reference field. Normalisation is tribal (“strip INV-”).
2	Key rule written. Not applied to portal vs email copies of the same invoice.
3	Documented key, normalisation, and collision behaviour. Workflow and ERP share the key.
4	Key rule versioned; tested against known resends; used before payment proposal.

Why it matters. Duplicate control and cycle-time clocks both die without a stable identity.

DS-02 — Master-data fitness

Question. Can AP trust vendor, entity, tax standing, and payment-term data enough to validate an invoice — and is AP forbidden from editing payment destination from the invoice?

Score	What it looks like in AP
0	Clerks create or edit vendors to “get it through.” Bank details copied from PDFs.
1	Steward exists. AP still applies one-off vendor fixes. Bank changes arrive as invoice footnotes.
2	Create/edit is ticketed. Bank-change path exists but is sometimes skipped under urgency.
3	Steward-only writes. Invoice-printed IBANs cannot update the master. Dual control on bank changes.
4	Steward-only writes, dual control evidenced, AP identity cannot write payment destination, and blocked/incomplete vendors fail closed.

Why it matters. The classic loss path is payment-master change, not OCR. Data readiness includes *who may write*, not just completeness percentages.

DS-03 — PO and GR readability

Question. Can a processor (or an agent at L0) read remaining PO quantity/value, GR quantity, and status *for the invoice’s entity and supplier* without a side spreadsheet?

Score	What it looks like in AP
0	Match is visual. Remaining value is phoned to the buyer.
1	ERP shows PO and GR. Interface lag is unknown. Plant receipts live in a warehouse system AP cannot see.
2	Reads work for stock POs. Service confirmations and blanket POs are opaque.

Score	What it looks like in AP
3	Documented objects for PO remaining, GR cumulative, status. Known interface lag is written (and treated as a fail-closed if stale).
4	Reads are API- or report-stable; heartbeat on GR/PO interfaces; stale feed stops match — it does not guess.

Why it matters. Matching on a stale GR extract fabricates “received.” That is not a model problem; it is a data-trust problem.

DS-04 — Evidence store

Question. When an invoice is parked, matched, or paid, can you retrieve the source file, the extract, the decision, and who overrode what — without asking a vendor to “pull a log”?

Score	What it looks like in AP
0	PDF lives in email. ERP has a document number. No join.
1	Archive exists. Search is by date. Override history is a comment.
2	Source attached in workflow. Agent or OCR traces live only at the vendor.
3	Buyer-controlled store: source, hash or equivalent, extract, decision, actor. Retention follows finance policy.
4	Store is the system of record for evidence; vendor traces are copies; agents cannot delete; export tested.

Why it matters. Permission to act requires proof. If proof evaporates with a vendor contract, you cannot promote an agent and you cannot defend a close.

DS-05 — Integration honesty

Question. Are capture, workflow, ERP, tax engine, and bank/payment file *distinct* systems with known owners — or a pile of CSVs and a shared inbox that “is the integration”?

Score	What it looks like in AP
0	People are the interface. CSV on desktops.
1	Some APIs or IDocs. Failures discovered when suppliers chase.
2	Interfaces listed. No heartbeat. Reruns can duplicate.
3	Interface register: source, target, owner, failure behaviour, at-most-once note for anything that creates a document.
4	Heartbeats, fail-closed on stale data, duplicate-safe reruns, owner on-call for payment week.

Why it matters. Agents amplify a bad interface. A retry that creates a second ERP invoice is a control incident, not an IT ticket.

DS-06 — Measurement extract

Question. Can you produce, from systems you control, last-quarter counts for invoices received, posted, exceptions by reason, cycle time, and open aged items — without a project?

Score	What it looks like in AP
0	Volume is "busy." No extract.
1	Vendor dashboard. Does not reconcile to ERP.
2	ERP counts exist. Exception mix and cycle time need a one-off query.
3	Repeatable extract: volume, mix, ageing, cycle-time clock defined (receipt → ready-for-payment or equivalent).
4	Repeatable extract, clock definition versioned, mix (channel × type × entity) available, used monthly.

Why it matters. The baseline KPI worksheet is empty without this. You cannot later claim improvement you cannot measure now.

Domain 2 evidence note. Artefacts: invoice-key rule, vendor/bank SoD extract, PO/GR data dictionary, archive pointer, interface register, last-quarter volume extract.

Domain 3 — Controls & audit

Would Internal Audit rather sample this operation — or write a finding the week an agent posts its first invoice?

CA-01 — Delegation of authority is a table

Question. Is DoA a current table (role, entity, amount, category) that workflow actually uses — not a PDF in a drawer?

Score	What it looks like in AP
0	Approvals are "whoever is around."
1	DoA PDF exists. Workflow limits are stale. Self-approval happens.
2	Table in workflow. Deputies are informal out-of-office texts.
3	Versioned DoA; workflow cites it; expired delegates fail; self-approval blocked.
4	Controller-owned table; change-controlled; sample of approvals vs table in the last 12 months.

Why it matters. An approval agent may route. It must not become the approver. If DoA is mush, routing only accelerates the wrong signature.

CA-02 — Segregation of duties on money

Question. Are invoice entry, vendor-bank change, and payment release incompatible duties — for people *and* for any service account already in use?

Score	What it looks like in AP
0	Same person can enter, change bank, and release. Or a shared login does.
1	Policy says SoD. Shared mailbox or "AP robot" account has combined rights.
2	Human SoD holds. Project/service accounts have not been reviewed.
3	Human and system identities reviewed. Incompatible pairs listed. Exceptions have expiry.

Score	What it looks like in AP
4	Quarterly access cert includes agent/service identities. Payment-release verb cannot sit on a capture identity.

Why it matters. Standing up an agent on a privileged token collapses SoD in an afternoon.

CA-03 — Duplicate control before cash

Question. Is exact-key duplicate checking performed before post *and* again before the payment proposal — and is a paid collision treated as an incident?

Score	What it looks like in AP
0	Clerks remember invoices. No systematic check.
1	ERP blocks exact same reference sometimes. Portal and email copies get through.
2	Check at entry. Payment run is not re-checked. Near-duplicates are ignored or guessed.
3	Exact-key rule documented; run at entry and at proposal; paid collision has an incident path. Near-duplicates are a <i>flag</i> , not a fraud label.
4	Both gates evidenced; false-negative sample includes paid population; near-duplicate review is owned.

Why it matters. This is not a fraud-detection claim. It is a payment-hygiene control. Missing it is a hard brake on execute-level agents.

CA-04 — Bank-change path

Question. Can a new IBAN or account on an invoice, email, or portal message *ever* update the payment destination without an official dual-control path?

Score	What it looks like in AP
0	Clerks update bank details from the PDF “to help the supplier.”
1	Policy forbids it. Urgency overrides it. No hold code.
2	Ticket path exists. Invoice-footer changes still happen.
3	Document-driven bank write is impossible in the system AP uses. Hold is mandatory. Out-of-band verification per Treasury policy.
4	Technical impossibility + hold + dual control + incident class if anyone acted on document text.

Why it matters. Payment destination is a human class at every autonomy level in Evidence Room. If the path is porous, do not add an agent that reads invoice text.

CA-05 — Override and hold hygiene

Question. When someone bypasses a match, a hold, or a workflow step, is it logged with a reason, a person, and a sample review — or is “force through” the unofficial SOP?

Score	What it looks like in AP
0	Force-post / lift-hold with no log.

Score	What it looks like in AP
1	Comments. No report.
2	System logs exist. Nobody reads them. Force-path is how month-end works.
3	Reason codes; named roles for force-path; weekly override rate visible to AP Manager.
4	Override sampled; force-path dual-controlled above a cap; rising override rate freezes any automation change.

Why it matters. An agent at L3 with a culture of force-path is an agent whose gates will be walked around.

CA-06 — Audit trail usable by a stranger

Question. Could Internal Audit, given an invoice ID, reconstruct source → extract → match/exception → approval → payment recommendation → release, with actors and timestamps, in one sitting?

Score	What it looks like in AP
0	No. Evidence is email.
1	ERP document flow only. Capture and approval live elsewhere with different keys.
2	Most steps findable. Overrides and vendor-portal copies are the gaps.
3	End-to-end join documented. A stranger can follow a sample of 10.
4	Join tested in the last audit cycle; gaps logged as findings with owners; agent actions (if any) already in the same trail.

Why it matters. If a human trail is incomplete, an agent trail will not magically appear.

Domain 3 evidence note. Artefacts: DoA version, SoD pair list + access extract, duplicate rule, bank-change SOP, override report, audit sample pack.

Domain 4 — Exception economics

Do you know what breaks, what it costs in time, who owns it, and whether the same break repeats — or is “exceptions” a single tired pile?

EE-01 — Volume and mix

Question. For the last complete quarter, can you state exception volume and the mix of the top codes (or your local reason codes)?

Score	What it looks like in AP
0	Unknown. The queue is “the queue.”
1	Headcount feel. “Mostly GR.” No count.
2	Total exception count. Mix is a pie chart from a vendor with different names.
3	Count and mix from <i>your</i> codes for a quarter, same scope (entity/channel).

Score	What it looks like in AP
4	Mix by code, owner, channel, entity; trend vs prior quarter; Misc under a cap.

Why it matters. First-wave agents should attack a named code family, not “AP.” Mix is how you pick.

EE-02 — Ageing is owned

Question. Are open exceptions aged by code and owner, with a named person for items past your own clock?

Score	What it looks like in AP
0	No ageing. Old items sit in personal folders.
1	Ageing report exists. Nobody is assigned. Month-end mass-closes occur.
2	Age buckets. Ownership is the last person who touched it.
3	Age × code × owner. Escalation clock written. Mass-close is a control event.
4	Ageing reviewed weekly; stalled paths get a second code (aged unresolved); write-off/reject above cap is human.

Why it matters. Agents that “clear ageing” by closing as no-issue are a control failure. You need the clock before you automate the chase.

EE-03 — Time on task

Question. Do you have *any* measured minutes (time study, timestamps, or sampling) for intake, match, and the top two exception types?

Score	What it looks like in AP
0	No. Cost is headcount divided by hope.
1	A vendor slide with “industry minutes.” Not your clocks.
2	One-off sample > 12 months old, happy path only.
3	Dated sample or system timestamps for intake + one high-volume exception, method written.
4	Repeatable timestamps; review/override minutes included; used in a cost model the Controller will look at.

Why it matters. Hours-released arithmetic without a time study is storytelling. This diagnostic will not invent a minute standard.

EE-04 — Repeat rate

Question. Do you know which supplier × exception-code (or buyer × code) pairs recur, and does anyone own a fix request?

Score	What it looks like in AP
0	Every invoice is new.

Score	What it looks like in AP
1	People “know” the bad suppliers. No list.
2	Informal watchlist. No fix owner.
3	Repeat report for a look-back you define. Top pairs have a named buyer or process owner.
4	Repeat report drives a fix request with evidence; recoding to hide repeats is sampled.

Why it matters. Root-cause work is how volume falls. An agent that chases the same missing GR 200 times is a very expensive printer.

EE-05 — Owner per break

Question. For missing PO, price variance, missing GR, duplicate, tax, and coding, is there a responsible *role* (not “AP will chase”)?

Score	What it looks like in AP
0	AP owns everything by default.
1	Unofficial: warehouse for GR, buyer for price — when they answer.
2	RACI poster. Live queue still dumps to AP.
3	Routing table: code → role → next action → clock. AP coordinates; does not absorb.
4	Table versioned; reject-to-sender when the owner is wrong; AP Manager reviews unowned aged items.

Why it matters. Exception Triage is useless without a routing table. Talent problems often hide here.

EE-06 — Unit cost honesty

Question. Do you have a Controller-acceptable view of what AP processing costs (people, vendors, allocations) per invoice or per exception — even if rough?

Score	What it looks like in AP
0	No cost model.
1	Benchmark from a conference. Not your ledger.
2	Headcount ÷ invoices. Ignores vendor, overtime, and plant AP.
3	Written cost model: components listed, volume milestone defined, Controller has seen it.
4	Versioned model used quarterly; inference/tool cost has a line if tools exist; no unsigned “savings” on the same page.

Why it matters. APQC’s published spread (\$0.38 vs \$0.92 per \$1,000 revenue to process AP) is *their* research, not your target. Your business case starts from *your* cost model. If you do not have one, the business-case starter stays at inputs-only.

Domain 4 evidence note. Artefacts: exception mix extract, ageing snapshot, time-study ID, repeat report, routing table, cost-model version.

Domain 5 — Talent & operating model

Is there a human operating system an agent could join — or would the agent join a hero culture?

TO-01 — Named roles, not heroes

Question. Are Head of AP, AP Manager, exception owners, payment preparer, payment authoriser, GR owner, buyer, and master-data steward identifiable people (or role pools) with deputies?

Score	What it looks like in AP
0	"The AP team." One person who knows everything.
1	Org chart. Deputies are implicit. Payment authoriser sometimes enters invoices.
2	Roles named. Two seats vacant; work absorbed silently.
3	Incumbents and deputies listed. SoD-sensitive seats are not combined "just this month."
4	Quarterly incumbency check; absence does not move accountability to a tool or a vendor.

Why it matters. Every Evidence Room agent requires a human Accountable. A vacant seat means the agent has no owner.

TO-02 — Capacity vs queue

Question. Is workload visible (invoices, exceptions, aged items per FTE or per queue) so you can see whether an agent would remove work or add review work onto an overloaded desk?

Score	What it looks like in AP
0	Overtime is the dashboard.
1	Volume known. Staffing is annual budget folklore.
2	Occasional workload snapshot. No link to exception mix.
3	Queue and FTE visible monthly; review time for any existing automation counted.
4	Capacity model; new agent review minutes estimated <i>before</i> go-live; AP Manager can refuse a pilot that adds net load.

Why it matters. L1 recommend increases human reading load. If the desk is already underwater, start smaller.

TO-03 — Skills to supervise, not only to process

Question. Does anyone in AP or Finance Systems have the skill to read a miss log, a false-positive report, and an override reason — or is "AI" owned by a vendor success manager?

Score	What it looks like in AP
0	No internal owner. Vendor runs the weekly call.
1	One enthusiast. No deputising. No QA method.
2	FinSys can pull logs. AP cannot interpret them. Or the reverse.

Score	What it looks like in AP
3	Named AP owner + FinSys owner. They can explain last month's errors to the Controller.
4	Supervisory skill in the operating rhythm (weekly miss/override review). Vendor is a supplier, not the Accountable.

Why it matters. Deloitte's 2026 survey found 63% of finance teams have fully deployed and actively use AI, while only 21% report clear, measurable ROI. Supervision skill is part of why those numbers diverge. That is their finding, not a promise that supervision produces ROI.

TO-04 — Supplier and internal chase is designed

Question. Are supplier and internal follow-ups templated, clocked, and sent from an official identity — or is chase personal email and chat?

Score	What it looks like in AP
0	Personal Outlook. Tone and facts vary.
1	Shared mailbox. No templates. No clock.
2	Templates exist. Not used under pressure. Directors get chased ad hoc.
3	Templates, clocks, allow-listed recipients, official identity. Commercial concessions forbidden in chase text.
4	Send log retained; empty "nudge" to hit a clock is a defect; legal mail is a different path.

Why it matters. Supplier Resolution and Internal Follow-Up agents draft; they should not invent a commercial position. If chase is already chaotic, drafts will scale the chaos.

TO-05 — Change and training path

Question. When match tolerance, reason codes, or intake channels change, is there a way processors actually hear about it?

Score	What it looks like in AP
0	Change is a hallway conversation.
1	Email blast. No confirmation.
2	SOP updated. Floor not re-trained.
3	Change record, communication, and a quiz or floor check for material changes.
4	Same path will be used for agent instruction changes. Processors know how to fall back to the no-agent SOP.

Why it matters. An agent instruction edit is a workflow change. If humans already miss SOP updates, they will miss autonomy changes.

TO-06 — Fallback without the tool

Question. If capture, workflow, or a future agent is down on payment week, can AP still receive, park, and prepare a payment proposal in the ERP you already use?

Score	What it looks like in AP
0	Work exists only in a vendor UI. Nobody remembers the fallback.
1	"We would figure it out."
2	SOP mentions fallback. Not drilled.
3	Fallback SOP; payment path works without the satellite tool; last drill dated.
4	Semiannual drill; minutes-to-disable a tool identity known; evidence store export exists.

Why it matters. Agents must not become the system of record. BCP is operating-model work, not an IT appendix.

Domain 5 evidence note. Artefacts: role/deputy list, workload snapshot, vendor vs internal RACI, chase templates, last change comms, fallback drill date.

Domain 6 — Agent readiness

Have you earned the right to give a tool a job — or are you scoring a slide?

AR-01 — Charter thinking

Question. For any existing bot, OCR, or "AI invoice" tool, is there a written job: in-scope, exclusions, owner, and what it must not do (pay, change bank, approve, invent GR)?

Score	What it looks like in AP
0	Tools have slogans. No job description.
1	Vendor scope statement. No exclusions. No human owner.
2	Internal owner. Exclusions implied. Payment and bank not explicitly forbidden.
3	One-pager: purpose, exclusions including the four human classes, owner, start-in-observe.
4	Charter-grade document; tools/entitlements listed; used in access reviews.

Why it matters. If today's tool has no job description, tomorrow's agent will inherit the same vagueness.

AR-02 — Observe-first discipline

Question. Have you ever run a new AP automation in shadow (observe / recommend only) against labelled outcomes before giving it a write?

Score	What it looks like in AP
0	Go-live is the test.
1	Vendor UAT with happy-path samples.
2	Pilot on live invoices with write access "limited." No labelled miss log.
3	Shadow or labelled historical pack before write. Precision and misses recorded.
4	

Score	What it looks like in AP
	Observe → recommend → prepare is policy. Skipping a level requires a recorded refusal of this diagnostic's intent — and should not happen.

Why it matters. Evidence Room default is L0 Observe. A culture that cannot shadow cannot promote.

AR-03 — Human classes understood

Question. Is it explicit, in writing, that payment release, vendor bank-change, policy exceptions, and legal disputes stay human — regardless of model confidence?

Score	What it looks like in AP
0	Not discussed. "If the AI is sure..." is heard in the corridor.
1	Payment release is human. The other three are assumed.
2	All four stated in a slide. Not in a control or SOP.
3	Written rule. Urgency does not move a case across the line. Confidence is not a second signature.
4	Written rule plus reason codes that force <code>human_required</code> . Sampled.

Why it matters. This is the cheapest control you can install today. It does not require software.

AR-04 — Kill-switch and identity

Question. Can a named human disable a capture bot, workflow automation, or service account *today* without opening a two-week ticket — and do you know which identity it uses?

Score	What it looks like in AP
0	Unknown identity. Vendor must be called.
1	Identity known. Disable path is theoretical.
2	FinSys can disable. AP Manager cannot trigger it. Untested.
3	Named kill-switch owners (AP Manager and Controls). Identity inventory exists.
4	Drill in the last six months; disable minutes recorded; scope can be stack / tool / vendor / entity.

Why it matters. L3 without a kill-switch is not managed autonomy. Even L0 tools should be disable-able.

AR-05 — Prompt and document distrust

Question. Is inbound invoice and email text treated as untrusted data — or have you already seen a document "instruction" change a process?

Score	What it looks like in AP
0	Clerks follow footer instructions (new bank, "please approve").
1	People know not to. No system help.
2	Policy says documents are not instructions. Not tested.

Score	What it looks like in AP
3	Bank and approval paths ignore document imperatives. Suspicious hidden text has a quality hold.
4	Explicit injection cases in any tool test pack; tools cannot fetch arbitrary URLs from invoices.

Why it matters. Agents that read PDFs will meet hostile or accidental instructions. The control is architectural, not “better prompting.”

AR-06 — Ambition matches evidence

Question. Is the current AP-AI ambition (none / experiment / one pilot / “autonomous AP”) consistent with your other five domain scores — in this sitting, not in a strategy deck?

Score	What it looks like in AP
0	Public commitment to autonomous AP or guaranteed savings. No measurement.
1	Multiple tools in flight. No single owner. Success = vendor slide.
2	One experiment. Success criteria unwritten.
3	Ambition is one bounded agent at observe/recommend. Success = miss log and evidence pack.
4	Written ambition: one slice, one owner, no execute until Measure. Board language matches this diagnostic.

Why it matters. Over-ambition is a control risk. Under-ambition is a delay. This question is the honesty check.

Domain 6 evidence note. Artefacts: any existing tool charter, shadow/UAT pack, human-class note, identity list, injection/quality hold example, written ambition statement.

Maturity model

Map **overall score** (average of 36 questions, 0.0–4.0) to one level. Use the **hard brakes** even if the average is higher.

Level	Name	Overall score	What this organisation is
1	Unprepared	0.0–0.7	Path, data, and controls are not stable enough to supervise a tool. Do not buy an agent. Fix identity, SoD, duplicate, and channel inventory.
2	Experimental	0.8–1.5	Some process language exists. Measurement is thin. A shadow on historical invoices is the most you should attempt — and only on a slice.
3	Governed pilot	1.6–2.3	You can name owners, codes, and a first agent. Observe and recommend are in play. Prepare is conditional. Execute is not earned.
4	Earned scale	2.4–3.1	Evidence packs exist. One increment at a time (level <i>or</i> scope). Human classes hold under pressure.
5	Institutional	3.2–4.0	

Level	Name	Overall score	What this organisation is
			Agents are jobs in an operating system: versioned instructions, sampled execute, kill-switch drilled, measurement owned by Finance — not a vendor. Still not “autonomous payment.”

Hard brakes (override the average).

Condition	Effect
Domain 3 (Controls) average < 2.0	Cap maturity at Experimental , regardless of overall.
CA-03 or CA-04 scored 0	Cap at Unprepared for any money-adjacent tool. Fix duplicate and bank-change first.
Domain 2 (Data) average < 1.5	Cap at Unprepared for live agents. Shadow on a static file is the ceiling.
More than 12 answers tagged G	Result is indicative . Do not use it to approve spend.
AR-06 scored 0	Record a governance finding. Ambition is a risk, not a score to “work on later.”

What to do. Read your level as a permission, not a grade.

How. Apply hard brakes after arithmetic.

Who. Challenge applies brakes. Scorer does not negotiate them away.

What can go wrong. Calling Level 3 “ready to scale” because a vendor uses that word.

Control. The table above is the only mapping. Do not add stars or traffic lights that contradict it.

Measure. Level, domain range, brakes applied (yes/no and which).

Evidence. This page, dated, with the six domain scores printed.

Opportunity heatmap

What to do. Pick up to five candidate opportunities (invoice slices or exception families). Score each on four axes using *your* data. Plot volume × pain on the page; use control risk and data readiness as filters, not as decoration.

How.

For each candidate, score 1–5 on each axis. Use the rubric. If you lack data, score the axis **1** and mark **unknown**.

Axis	1	3	5
Volume	Rare in your last quarter	Material weekly queue	Dominant share of invoices or exceptions in the chosen scope

Axis	1	3	5
Pain	Low age, low supplier noise, low overtime	Visible stall or rework	Aged, escalated, or consuming disproportionate minutes (from <i>your</i> EE-03, not a benchmark)
Control risk	Break cannot move cash or identity	Operational error (wrong park, slow chase)	Touches payment destination, duplicate pay, tax position, GR invention, or DoA
Data readiness	Cannot read the objects needed	Reads exist; stale or incomplete	Stable keys, readable PO/GR/vendor, evidence store for this slice

Plot. X = Volume (1–5). Y = Pain (1–5). Each bubble is a candidate.

Filter (do not skip).

If	Then
Control risk = 4 or 5	Not a first-wave <i>execute</i> candidate. Eligible for observe/recommend only, or for a human-class packaging agent (e.g. Duplicate flag, Payment review later).
Data readiness ≤ 2	Do not agentise. Fix the extract or interface.
Volume ≤ 2 and Pain ≤ 2	Park. It is a pet project.
Volume ≥ 4, Pain ≥ 4, Control risk ≤ 3, Data readiness ≥ 4	First-wave shortlist. Still starts at L0.

Who. AP Manager proposes candidates from EE-01 mix. Control owner scores control risk. Finance Systems scores data readiness. Head of AP picks **one** first slice.

What can go wrong. Plotting “AI matching” as a single bubble. Matching is a job; missing GR is a code family; they are different bubbles. Plotting vendor-promised touchless rates (Hackett’s 60% average is *their software-adopter figure*, not your plot).

Control. Every bubble cites a volume source (extract ID or **unknown**). Unknown volume cannot be scored above 2.

Measure. Count of bubbles that survive the filter. If zero, the next step is data or control work, not a Starter Kit implementation.

Evidence. Heatmap table in this file or the workbook.

Heatmap worksheet (buyer fills)

Scope for this plot: entity [] · period [] · extract ID []

ID	Candidate (slice or code family)	Vol 1–5	Pain 1–5	Ctrl 1–5	Data 1–5	Survive filter? Y/N	First agent (if any)
H1							
H2							
H3							

ID	Candidate (slice or code family)	Vol 1–5	Pain 1–5	Ctrl 1–5	Data 1–5	Survive filter? Y/N	First agent (if any)
H4							
H5							

ACME — ILLUSTRATIVE. Entity Plant North, Q1. H1 Missing GR on stock POs: Vol 5, Pain 5, Ctrl 3 (GR invention risk noted), Data 4 → survives as Goods Receipt *observe*. H2 "Autonomous payment": Vol 5, Pain 3, Ctrl 5, Data 3 → fails filter. H3 Email PDF intake identity: Vol 4, Pain 3, Ctrl 2, Data 4 → survives as Intake L0.

Ten-agent opportunity overview (first wave)

The Professional library is sixteen agents. **Do not stand up sixteen.** This diagnostic uses the first-wave ten: the capture–qualify–resolve path plus the orchestrator. Agents 06 PO Quality, 11 Statement, 12 Payment Proposal Review, 13 Close, 14 Reporting, and 15 Root Cause wait until Measure exists.

#	Agent	Job in one line	Typical first owner	Earliest honest level	Heatmap fit
01	Invoice Intake	Capture, classify, extract, park with an evidence pack	AP Supervisor / Intake Lead	L0	High volume, low–med control risk if it cannot write bank or post
02	Invoice Validation	Prove complete, coded, master-data-fit before match	AP Team Lead	L0	After Intake evidence exists
03	Matching	2-way / 3-way / tolerance; state the residual	AP Processor (PO desk)	L0	Only if PI-04 ≥ 2 and DS-03 ≥ 3
04	Exception Triage	Code, owner, next action, SLA	AP Exception Lead	L0	Only if PI-05 ≥ 2 and EE-05 ≥ 3
05	Goods Receipt	Find missing/partial GR evidence; never invent a receipt	Warehouse / Plant AP liaison	L0	Common high-volume bubble; execute is rare
07	Approval	Build the packet; route inside published DoA	AP Supervisor	L0	Only if CA-01 ≥ 3; agent is never the approver
08	Supplier Resolution	Draft the supplier question; send only if a send-gate allows	AP Vendor Desk	L0	After triage codes are clean
09	Internal Follow-Up	Chase receiver, buyer, cost-centre	AP Query Desk	L0	After routing table exists
10		Flag exact and near-duplicate and unusual		L0	

#	Agent	Job in one line	Typical first owner	Earliest honest level	Heatmap fit
	Duplicate & Anomaly	patterns — not a fraud finding	AP Controls Analyst		Pair with Intake from day one
16	Orchestrator	Sequence, gate, stop; hold the queue	AP Manager	L0	Required as soon as two specialists exist

What to do. From surviving heatmap bubbles, pick **one** specialist plus Intake (01) and Duplicate (10) at L0 if those are not already the specialist. Add Orchestrator (16) the week a second agent appears.

How. Use the Starter Kit blueprints. Do not write prompts this afternoon.

Who. Head of AP picks. Control owner confirms the earliest level.

What can go wrong. Starting with Payment Proposal Review or “the orchestrator that posts.” Starting without Duplicate when Intake is live.

Control. First-wave list is closed. Adding Agent 12 to wave one requires CA-03 ≥ 3 , CA-04 ≥ 3 , and a Treasurer in the room — and it still starts at L1 recommend only.

Measure. Number of first-wave agents proposed (target: 1 specialist + required companions, not 10).

Evidence. Named first agent, slice, owner, on the interpretation page.

Baseline KPI worksheet

What to do. Fill what you can from the last complete quarter (or your last complete month, if that is all you have). Leave blanks. Do **not** paste conference benchmarks into the target column.

How. One scope line. One clock definition. All figures from *your* systems.

Scope

Field	Buyer value
Legal entity / company code	[]
Period	[]
Channels included	[]
Invoice types included	[]
Clock start (receipt at first AP channel / invoice date — pick one)	[]
Clock end (posted / ready-for-payment — pick one)	[]
Extract ID or report name	[]
Prepared by / date	[]

Counts (activity — not success)

Metric	Formula	Your figure	Source	Unknown?
Invoices received	Count at clock start			
Invoices posted or ready-for-payment	Count at clock end milestone			
PO invoices	Count			
Non-PO invoices	Count			
Credits	Count			
Exceptions opened	Count of non-STP items			

Operational (leave blank if you cannot compute)

Metric	Formula	Your figure	n (denominator)	Source
First-pass / STP (your definition — write it)	STP invoices / invoices in scope			
Write your STP definition	[]	—	—	—
Median cycle time	Median (clock end – clock start)			
p90 cycle time	90th percentile of the same			
Open exceptions now	Count			
Aged exceptions (your bucket, e.g. > [] days)	Count and value			
Top exception code	Name + count			
Second exception code	Name + count			
Missing GR (if you can identify it)	Count and value			
Exact duplicates flagged	Count (not a savings figure)			
Paid duplicate collisions known	Count (incident)			

Cost (only if a model exists)

Metric	Formula	Your figure	Model version	Controller seen? Y/N
AP processing cost in period	Sum of listed components			
List components	[people / vendor / allocation / other]	—	—	—
Cost per invoice				

Metric	Formula	Your figure	Model version	Controller seen? Y/N
	AP cost / invoices at stated milestone			

Do not fill

- Industry target STP
- “Expected AI savings”
- Hackett touchless % as a goal
- APQC \$ per \$1,000 revenue as *your* target

Those sources may appear in a narrative appendix as market context. They are not rows on this worksheet.

Who. Analyst pulls. AP Manager defines the clock. Controller initials the cost row or it stays blank.

What can go wrong. Mixing entities. Changing the clock later and calling the delta “improvement.” Putting a target in the figure column.

Control. Blanks allowed. Invented figures not allowed. Period and scope printed on every later scorecard that claims a baseline.

Measure. Count of completed numeric rows vs 18 possible. Completeness is a Data & systems signal.

Evidence. This worksheet, frozen, ID [BUYER]/Evidence/Diagnostic/Baseline_<period>.

Business-case starter (inputs only)

What to do. Structure a conversation the CFO can have without a fabricated ROI. Capture inputs. Build three *scenarios of assumption*, labelled as such. Do not multiply them into a promised saving in this diagnostic.

How. Complete the input table. Then write conservative / base / upside as **assumption sets**, each with a named owner who would have to validate any later cash line.

Inputs (facts or explicit unknowns)

Input	Value	Source	Unknown
Baseline invoices / period (from worksheet)			
Baseline exceptions / period			
Measured minutes — intake (if any)		Time study ID:	
Measured minutes — top exception (if any)		Time study ID:	

Input	Value	Source	Unknown
Loaded cost per AP hour (if Controller agrees to use one)			
Current vendor / capture spend in period			
Proposed tool / project cash out (if known)		Quote ID:	
Proposed internal hours to implement (estimate — label ESTIMATE)			
Slice proposed (entity × channel × type)		Heatmap ID:	
First agent			
Autonomy ceiling for the case period	L0 / L1 / L2 (L3+ does not belong in a Tier 0 case)		

Scenario structure (no computed promise)

Write qualitative and, only where an input exists, arithmetic *candidates*. Every cash or hour figure remains **UNVALIDATED** until the Controller signs a later pack.

	Conservative	Base	Upside
Scope	Narrower than the heatmap bubble (e.g. one vendor set)	The surviving bubble	Bubble + one related code family
Autonomy	L0 only for the case period	L0 → L1 if Measure exists	L1 → L2 prepare only
Volume assumption	Lowest plausible in-scope count you can defend from the extract	Worksheet count	Worksheet count + stated mix shift — do not assume mix gets easier
Minute assumption	No minute reduction claimed if EE-03 < 3	Reduction only on the <i>measured</i> task, including new review minutes	Same, with a stated range — still ESTIMATE
Error assumption	Miss rate no better than today; possible new false positives	Miss rate unknown; review load included	Still not “zero errors”
Cash line	Tool cost only (cash out)	Tool cost + ESTIMATE hours	Same; no duplicate-recovery cash unless you already have proposal-diff evidence
What would have to be true	Shadow pack exists	Owner staffed; codes clean	Controller willing to look at hours
What would falsify it	Cannot extract baseline	Override becomes the path	Rework or FNR worsens
Owner of the assumption	AP Manager	Head of AP	Controller (cash) / Head of AP (hours)

Forbidden in this starter

- A single ROI percentage.
- Annualising a two-week pilot.
- Counting flagged duplicates as saved cash.
- Using APQC, Deloitte, or Hackett figures as *your* upside.

Who. Head of AP drafts. Controller refuses unsigned cash. CFO reads the falsifiers.

What can go wrong. A spreadsheet that multiplies “industry 60% touchless” by your volume. That number is Hackett’s software-adopter research, not an input you measured.

Control. This section contains no “expected saving” cell. If someone adds one, it is not this diagnostic.

Measure. Count of unknowns in the input table. More than half unknown → case is premature.

Evidence. Completed input table + scenario narrative, stored with the score.

ACME — ILLUSTRATIVE. Inputs: 2,100 Plant North email PDFs in a quarter; missing-GR minutes from a 40-invoice sample; no loaded rate signed. Conservative: Intake + Duplicate at L0, cash out = existing capture licence only. Base: L1 recommend on extract fields. Upside: L2 park-drafts. No savings line. Controller initials “no cash claim.”

How to interpret the score

What to do. Produce a one-page read for the CFO. Use the pattern below. Do not average your way to comfort.

How — the one-page script.

1. **Overall.** “We scored **X.X / 4.0 (N / 144)**. Maturity **level name**. Brakes applied: [none / list]. Guess-tagged answers: [n].”
2. **Shape.** Print six domain scores. Name the lowest domain in one sentence: what is missing in AP terms.
3. **Permission.**
 - Level 1 → no agent spend. Fix CA-03/CA-04/DS identity as applicable.
 - Level 2 → historical or shadow only, one slice.
 - Level 3 → Starter Kit; one agent; L0/L1.
 - Level 4 → Professional operating system; still one increment.
 - Level 5 → operate; do not relax human classes.
4. **First slice.** The single heatmap survivor you will actually staff. If none, say so.
5. **Not a finding of value.** The score does not say agents will pay for themselves. Deloitte’s 21% “clear, measurable ROI” figure is a reminder that deployment ≠ return. Cite it only as context.

Decision table

If you see	Say	Do next
Controls < 2.0		

If you see	Say	Do next
	"We are not allowed to execute. We may not even be allowed to experiment on live invoices."	Bank-change + duplicate + SoD work
Data < 1.5	"We cannot measure or feed an agent honestly."	Channel volume + invoice key + extract
Exceptions strong, process weak	"We know the pile; we do not know the path."	As-done map before Triage
Agent readiness >> other domains	"Ambition is ahead of the shop."	Freeze new tool spend; run this diagnostic's hard brakes
All domains ~2 with few guesses	"Governed pilot is honest."	Buy or use the Starter Kit; book a Monday workshop
High score, 15 guesses	"We performed confidence, not diagnosis."	Two-hour evidence pull; re-score

Who. Head of AP writes the page. Challenge signs "brakes correctly applied." CFO receives one page plus the baseline worksheet.

What can go wrong. Sending the 36 rubrics without interpretation. Sending interpretation without domain scores.

Control. One page. Six numbers. One slice. One brake statement.

Measure. Whether the CFO can repeat the permission sentence without notes.

Evidence. The one-pager, dated.

What this is NOT

This diagnostic is **not**:

- A guarantee of savings, ROI, or payback.
- A fraud-detection assessment or a statement that duplicates will be caught.
- A regulatory, tax, SOX, or statutory-compliance opinion.
- An accounting-accuracy certificate or an audit.
- Permission to release payments, change vendor bank data, grant policy exceptions, or settle legal disputes by model.
- A ranking of ERPs or vendors.
- A substitute for the Starter Kit (Tier 1) or the Professional operating system (Tier 2).
- Valid if weights were changed, if industry targets were typed into the baseline, or if one person scored without a challenge.

Next step

If maturity (after brakes)	Product	What you do this week
1 Unprepared	Stay on Tier 0	Fix the named brake (usually duplicate, bank-change, invoice identity, or SoD). Re-score when artefacts exist. Do not buy an agent.
2 Experimental	Optional: Starter Kit	Read the human-vs-agent page. Run a labelled historical sample on one slice. No live write.
3 Governed pilot	Starter Kit (US\$79)	Friday buy → Monday workshop. One bounded agent. L0.
4-5	Starter Kit then Professional	You still start one agent. Professional adds the full 16-agent library, control matrix, autonomy policy, and measurement dictionary.

Starter Kit gives an AP Manager a workshop they can run: operating model, ten blueprints, 27 exception codes, job-description and instruction templates, twelve KPIs, a 90-day path for one agent.

Professional is the operating system: sixteen charters, full taxonomy cards, governance, process-mapping method, KPI dictionary. Buy it when you have a named owner and a frozen baseline — not as a substitute for either.

Proof before permission. Agents earn responsibility. Evidence decides.

End of Evidence_Room_AP_AI_Readiness_Diagnostic.md

Evidence Room — AP AI Readiness Scorecard Instructions

Product: Evidence Room AP Agent OS — Tier 0 (Free Diagnostic)
Companion to: Evidence_Room_AP_AI_Readiness_Diagnostic.md
Workbook: Evidence_Room_AP_AI_Readiness_Scorecard.xlsx (generated later; this file is the contract for that workbook)
Brand idea: Proof before permission.
Line: Agents earn responsibility. Evidence decides.
Version: 1.0
Classification: Operating instructions for the scorecard. Not a scoring model of its own. If this file and the diagnostic disagree on a formula, the diagnostic wins.

Purpose

The markdown diagnostic is the method. The Excel scorecard is the **calculator and evidence packet** a CFO can open without reading 36 rubrics in prose.

This file tells you how to build, fill, check, and interpret that workbook — including if you are still on a printed score sheet or a blank workbook because the `.xlsx` has not been generated yet.

What to do. Capture 36 scores, six domain averages, overall points, maturity, hard brakes, heatmap, baseline KPIs, and business-case inputs in one file.

How. Follow the sheet list and the cell rules below. Do not add weights, targets, or a composite “AI index.”

Who. Scorer (Head of AP or AP Manager) enters scores. Challenge (Controller or Control owner) locks the file after review. Analyst pastes extracts into the Baseline sheet only.

What can go wrong. A helpful analyst adds `IF` logic that lifts a 2 to a 3 when “AI project exists.” That file is no longer this diagnostic.

Control. Every computed cell is a `SUM` or `AVERAGE` of integers you typed, or a `VLOOKUP` / `IFS` against the maturity table copied verbatim from the diagnostic. No other scoring logic.

Measure. Lint checklist at the end of this file. A workbook that fails lint is not issuable.

Evidence. Dated workbook or printed sheets stored under `[BUYER]/Evidence/Diagnostic/`.

If the Excel file does not exist yet

What to do. Recreate the workbook from this specification in 20 minutes, or print the tables and compute with a calculator.

How — minimum viable packet (no Excel).

1. Print or copy the 36-question answer grid (Sheet `Scores` layout below).
2. Write 0–4 and a one-line note per question.
3. Average each domain of six; average all 36.
4. Apply hard brakes by hand using the diagnostic table.
5. Fill Baseline and Heatmap on paper.

How — blank workbook (Excel / Sheets / LibreOffice).

Create sheets in this **exact order and name**. Do not rename; Professional-tier packs and later automation will look for these names.

#	Sheet name	Purpose
1	<code>00_Cover</code>	Identity, date, initials, maturity result
2	<code>01_Read_me</code>	Short instructions + link to this file
3	<code>02_Respondent</code>	People, scope, ERP names
4	<code>03_Scores</code>	36 questions, 0–4, notes, tags
5	<code>04_Domains</code>	Six averages + range
6	<code>05_Maturity</code>	Overall, points, level, brakes

#	Sheet name	Purpose
7	06_Heatmap	Five candidate rows + filter
8	07_Agents	First-wave ten + your pick
9	08_BaseLine	KPI blanks
10	09_Case	Business-case inputs + scenario text
11	10_Interpret	One-page CFO script
12	11_Lint	Issuance checklist

Protect all cells except input ranges (listed per sheet). Colour **input cells** only. Do not colour scores green/red except the brake flag.

Who. Anyone may build the blank. Only the scorer and challenge may fill 03_Scores.

What can go wrong. Merging 03_Scores and 05_Maturity so a CFO edits a score while reading the result.

Control. Cover sheet states diagnostic version 1.0 and “weights: none.”

Measure. Sheet count = 12. Extra dashboard sheets fail lint.

Evidence. File name: Evidence_Room_AP_AI_Readiness_Scorecard_<ENTITY>_<YYYY-MM-DD>.xlsx.

How to run a scoring sitting (with or without Excel)

What to do. Complete the diagnostic in 30–40 minutes as specified in the markdown file.

How.

1. Open 00_Cover. Fill entity, period of *judgement* (today's date), and names. Leave the maturity cells empty — they calculate later.
2. Complete 02_Respondent before any scores. If you cannot name the payment authoriser and the bank-change path owner, stop and fetch names. Several control questions are unanswerable without them.
3. Score 03_Scores top to bottom. One integer per row. Tag E / 0 / G.
4. Do not look at 05_Maturity until all 36 are filled. Looking early produces score-chasing.
5. Challenge reviews any 3 or 4 that lacks an artefact pointer in the note.
6. Analyst (or scorer) fills 08_BaseLine from extracts. Blanks stay blanks.
7. AP Manager fills 06_Heatmap and 07_Agents after scores, not before.
8. Head of AP writes 10_Interpret using the script. Challenge types the brake statement.
9. Run 11_Lint. If it fails, do not send to the CFO.
10. Save, hash or write the filename + timestamp on the cover, store the file.

Who. See diagnostic “How to complete.”

What can go wrong. Filling heatmap first ("we already know we want matching"). Completing Baseline with Hackett or APQC numbers in the *figure* column.

Control. 08_Baseline figure cells must not contain the strings 0.38 , 0.92 , 60% , 63% , 21% , 14% , 3.5 , 59% unless you have independently measured the same thing in *your* ledger and the note says so. Those are published research figures, not your baseline.

Measure. Elapsed minutes (cover field). Target sitting: 30–40.

Evidence. Cover initials of scorer and challenge.

Sheet-by-sheet contract

00_Cover

Field	Rule
Product	Hard-coded: Evidence Room AP Agent OS — Tier 0
Diagnostic version	1.0
Buyer legal entity	Input
Sitting date	Input
Scorer name / role	Input
Challenge name / role	Input
Guess count	=COUNTIF('03_Scores'!Tag,"G")
Overall score	From 05_Maturity (link, not retyped)
Overall points	From 05_Maturity
Maturity level	From 05_Maturity
Brakes applied	From 05_Maturity
Indicative only?	TRUE if guess count > 12
File ID	Input after save

No charts on the cover. No vendor logo other than Evidence Room wordmark if present.

01_Read_me

Paste, short:

- Scores are integers 0–4. Between two scores, take the lower.
- No weights.
- Hard brakes can cap maturity below the average.
- Industry research stays off the Baseline figure column.
- This workbook does not promise savings, fraud detection, compliance, accounting accuracy, payment safety, or ROI.

Link to the markdown diagnostic for rubrics. The workbook should carry **question text + short rubric** so it is usable offline; full “why it matters” stays in markdown if space is tight.

02_Respondent

Block	Fields
Organisation	Legal name, entity/company codes in scope, AP location (SSC / plant / outsource)
Systems	ERP, capture, workflow, archive, tax engine, bank/payment file — names only
People	Head of AP, AP Manager, Controller, Control owner, Payment authoriser, Bank-change owner, FinSys, Internal Audit (informed)
Volume hint	Last complete quarter invoice count if known — else unknown (this is not the Baseline sheet)
Constraints	Recording/privacy notes if a later workshop will observe the floor

What can go wrong. Writing “SAP” when the live path is SAP + a plant warehouse system + a desktop register. List satellites.

03_Scores

One row per question. Do not split a question across rows.

Column	Name	Rule
A	ID	PI-01 ... AR-06 — locked
B	Domain	Locked
C	Question	Locked (from diagnostic)
D	Score	Integer 0–4 only. Data validation. Blank = incomplete, not zero.
E	Tag	E / 0 / G / blank
F	Note	≤ 200 characters. Artefact name or “none.”
G	Disagreement	Optional second score from challenge if they differed ≥ 2. Used score is still the lower of D and G.
H	Final	=MIN(D, G) if G filled, else D

Domain totals use **Final**, not the scorer’s first number.

Locked rubric helper (optional columns I–M). Five short cells: “0 looks like...”, “...4 looks like...”. If included, copy from the diagnostic; do not rewrite.

Validation.

- Score outside 0–4 rejected.
- Final blank if Score blank — do not coerce to 0. A zero must be typed.
- Completeness: 36 Finals filled before 05_Maturity displays a level (else “INCOMPLETE”).

04_Domains

Domain	IDs	Score formula
1 Process integrity	PI-01-PI-06	Average of six Finals
2 Data & systems	DS-01-DS-06	Average
3 Controls & audit	CA-01-CA-06	Average
4 Exception economics	EE-01-EE-06	Average
5 Talent & operating model	TO-01-TO-06	Average
6 Agent readiness	AR-01-AR-06	Average

Also compute:

- Domain range = max – min of the six averages.
- Flag if range ≥ 2.0 : **UNEVEN**.

Display all six. Do not chart a radar that becomes the CFO slide without the brake text.

05_Maturity

Cell	Formula / rule
Points	SUM of 36 Finals. Max 144.
Overall	Points / 36, one decimal.
Raw level	IFS on Overall: 0.0–0.7 → 1; 0.8–1.5 → 2; 1.6–2.3 → 3; 2.4–3.1 → 4; 3.2–4.0 → 5
Brake: Controls < 2.0	If Domain 3 < 2.0 then cap level at 2
Brake: CA-03=0 or CA-04=0	Cap level at 1 for money-adjacent permission (display "UNPREPARED for money-adjacent tools")
Brake: Data < 1.5	Cap level at 1 for live agents
Brake: Guess > 12	Set INDICATIVE = TRUE ; do not change the numeric level, but Cover shows "indicative — do not approve spend"
Brake: AR-06 = 0	Finding flag AMBITION_RISK
Issued level	Minimum of raw level and brake caps

Label cells (locked text): Unprepared / Experimental / Governed pilot / Earned scale / Institutional.

Do **not** convert points to a percentage of 144 for the cover. Points and 0–4 average are enough. A percentage invites exam-grade behaviour.

06_Heatmap

Five rows H1–H5.

Column	Rule
Candidate	Slice or exception family — not “AI”
Vol, Pain, Ctrl, Data	Integers 1–5. If unknown, 1 and tag unknown
Volume source	Extract ID or unknown
Survive?	Y only if Vol $\geq$ 4 AND Pain $\geq$ 4 AND Ctrl $\leq$ 3 AND Data $\geq$ 4 — or leave as a manual Y/N the AP Manager types after reading the diagnostic filter (preferred: manual , because first-wave observe on a high-control-risk bubble can be valid). Default formula helper: fail if Data $\leq$ 2 or (Vol $\leq$ 2 and Pain $\leq$ 2). AP Manager overrides with a note.
First agent	From the ten, or blank

No automatic ranking of agents.

07_Agents

Locked table of the ten first-wave agents (copy from the diagnostic). One input row:

Field	Input
First specialist	Dropdown of the ten
Companions	Default suggest 01 + 10 if specialist is not already 01 or 10; 16 if two specialists
Slice	Entity × channel × type
Ceiling this quarter	L0 / L1 / L2
Owner	Named human

08_Baseline

Copy the diagnostic worksheet tables. All figure cells are **inputs**. No default numbers.

Header: scope, period, clock start, clock end, extract ID.

Lint rule: figure cells empty is OK. Figure cells that equal published research constants fail lint (see strings above).

STP definition is a text cell. Without it, the STP figure is invalid even if numeric.

09_Case

Input table from the diagnostic. Three text boxes: Conservative / Base / Upside. No calculated “savings” cell. If a later builder adds `=hours*rate*volume`, that cell must be labelled **UNVALIDATED** and must **not** flow to Cover.

10_Interpret

Structured fields, 400 characters each:

1. Overall sentence (can be formula-assisted from Cover).
2. Lowest domain sentence (input).
3. Permission sentence (input; must match issued level).

- 4. First slice (link to 07_Agents).
- 5. What we are not claiming (locked boilerplate).

11_Lint

Yes/No checks. All must be Yes to issue.

Check	Pass if
L1	36 Final scores filled
L2	Every score 0–4 integer
L3	Domain scores displayed (6)
L4	Issued level uses brakes
L5	Guess count displayed
L6	If Guess > 12 then Cover says INDICATIVE
L7	Baseline has scope + period + clock definitions or explicit unknown
L8	Baseline figures do not contain forbidden research constants unless noted as buyer-measured
L9	No savings/ROI cell on Cover
L10	Scorer and challenge named
L11	Heatmap filled only after scores (sitting note)
L12	Diagnostic version 1.0 cited
L13	File stored with ID on Cover

Scoring hygiene

What to do. Keep the score honest under social pressure.

How.

Pressure	Response
"We just bought capture AI — that is a 4 on Agent readiness."	AR-01 is about a <i>charter</i> , not a licence. Score the paper.
"Audit passed last year."	CA-06 is whether a stranger can reconstruct <i>this</i> invoice path, not whether the opinion was clean.
"We're a top-quartile AP shop."	Not a question on the sheet. APQC's cost spread is not a self-score.
"Put 3s so we can buy the kit."	Starter Kit is useful at Level 2–3. Inflating to 4 wastes the Controller's time.
"Average us to 2.4 so we look Earned scale."	Brakes and unweighted average. No.

- Who.** Challenge owns hygiene. Scorer owns completeness.
- What can go wrong.** Scoring in a room with the vendor present.
- Control.** Vendor and implementation partner are not in the sitting. They may receive the issued one-pager later, not the raw 36.
- Measure.** Count of scores changed after first look at 05_Maturity . Any change needs a note. More than three changes after seeing maturity → re-sit.
- Evidence.** If Excel: a copy saved at “scores complete, maturity hidden” if you want a clean audit. Optional, not required.

How to read the issued workbook (CFO, 5 minutes)

- 00_Cover — level, points, brakes, indicative flag.
 - 04_Domains — which domain is the hole.
 - 10_Interpret — permission and first slice.
 - 08_Baseline — how empty is measurement.
 - Stop. Do not mine 03_Scores unless you are the challenge preparing an audit query.
- If Cover shows INDICATIVE, the meeting outcome is “get the extracts,” not “approve a platform.”

What can go wrong with the file itself

Failure	Control
Hidden weights in 05_Maturity	This contract; lint by a second person who reads the formulas
Traffic-light KPI sheet added	Remove before issue
Multiple entities averaged in one sitting	One workbook per entity; group roll-up is a later Professional product
Printed scores diverge from Excel	Cover file ID; paper marked “superseded”
Cloud sheet shared with edit rights to the whole AP team	Challenge locks; viewers only after issue

Measure (of the scorecard process)

Check	Definition
Sitting time	Minutes, cover field
Completeness	Finals filled / 36

Check	Definition
Guess rate	G / 36
Lint pass	All L1–L13 Yes
Post-maturity edits	Count

Evidence retention

Keep the issued workbook for the period your finance policy keeps management papers. This product does not set a statutory period. Agents (if any later) must not be able to edit or delete this file.

Relationship to later products

Product	Uses this workbook as
Starter Kit	Entry ticket: first slice and ceiling
Professional	Baseline ID cited on Measure packs
A future generated .xlsx	Must implement this contract. Pretty formatting is optional. Formula changes are not.

Disclaimers

- Completing the scorecard does not implement a control.
- A high score is not permission to release payments or to claim ROI.
- Research citations belong in narrative, not in Baseline figures.
- Evidence Room does not receive a copy of your scores unless you choose to send them.

Proof before permission. Agents earn responsibility. Evidence decides.

End of SCORECARD_INSTRUCTIONS.md