

EVIDENCE ROOM

AP Agent Starter Kit

Operating model, top agents, taxonomy, and first-90-days path

Tier 1 · US\$79

Operating Evidence · Not legal, tax, or accounting advice

AP Agent Starter Kit — Guide

Product: Evidence Room — AP Agent Starter Kit **Price:** US\$79 **Document ID:** ER-START-GUIDE-001
Version: 1.0 **Brand idea:** Operating Evidence **Licence:** Individual (see §14) **Audience:** Head of AP, Controller, Finance Transformation lead designing the first controlled agent pilots

Cover

Evidence Room AP Agent OS · Starter

AP Agent Starter Kit

Charter the first agents. Keep payment human. Earn responsibility.

This kit is the operating layer between “we should use AI in AP” and a controlled pilot. It is not software, not an ERP module, and not a promise of autonomous payables.

You will leave with: an operating model sketch, ten usable agent blueprints, a condensed exception taxonomy, decision frameworks, templates, a starter KPI scorecard, governance checklist, and a 90-day path.

START HERE — 12-step path

Complete in order. Do not skip to “prompting an agent” before steps 1–6.

1	Read §1 Operating Model (30 min)	Shared vocabulary with sponsor	■
2	Run Free Diagnostic or transfer scores (§5)	Band + dimension profile	■
3	Inventory channels, volumes, top exception pains	One-page scope note	■
4	Install condensed exception taxonomy (§4)	Code list in tracker	■
5	Baseline KPIs (§10) — 2 weeks if data thin	Baseline sheet	■
6	Map process for pilot slice (§7)	Process map (template)	■
7	Choose ≤3 agents from Top 10 (§3)	Pilot backlog	■
8	Write Job Descriptions + Instructions (§8–9)	Charters v0.9	■
9	Apply Human-vs-Agent framework (§6); set L0/L1 ceiling	Autonomy registry row	■
10	Complete Governance checklist (§11)	Sign-offs	■

11	Shadow 2–4 weeks; sample QA	Shadow log	■
12	Decide: demote / hold / promote to L1 recommend-only	Steering note	■

Hard stops throughout: no payment release by agent; no silent vendor bank changes; no inventing DOA; no expansion without scorecard evidence.

1. AP Agent Operating Model overview

1.1 What “Operating Evidence” means here

An AP agent is allowed to act only when Finance can show:

- 1 **Charter** — job, inputs, exclusions, owner
- 2 **Control** — what must never happen; escalation
- 3 **Exception logic** — coded failures, not inbox folklore
- 4 **Measurement** — KPIs with baseline
- 5 **Responsibility level** — L0–L4 earned, never assumed

If any of the five is missing, the agent stays at observe/recommend.

1.2 The stack (sixteen agents — Starter focuses on ten)

A01	Invoice Intake	Blueprint
A02	Invoice Validation	Blueprint
A03	Matching	Blueprint
A04	Exception Triage	Blueprint
A05	Goods Receipt	Blueprint
A06	PO Quality	Mentioned; Pro deepens
A07	Approval	Blueprint (routing assist only)
A08	Supplier Resolution	Blueprint
A09	Internal Follow-Up	Blueprint
A10	Duplicate & Anomaly	Blueprint
A11	Vendor Statement Reconciliation	Blueprint
A12	Payment Proposal Review	Blueprint (prepare only)
A13–A16	Close, Reporting, Root Cause, Orchestrator	Roadmap sketch; full in Professional

1.3 Responsibility ladder (mandatory)

L0	Observe	Score, log, shadow	Run live process
L1	Recommend	Drafts, suggestions	Decide / send / post packs,

L2	Prepare	Write drafts into systems on clear cases	Approve material actions; sample QA
L3	Execute within guardrails	Bounded execution	Own escalations; monitor
L4	Managed autonomy	Bounded domain, exception-only touch	Policy + assurance; never payment auth

Starter recommendation: Cap new programmes at **L1** for 60–90 days. L2 only after Professional-grade controls and scorecards.

1.4 Payment boundary (non-negotiable)

Agents may **prepare** payment proposals and **flag** holds. Humans **authorise** and **release** payments. No Starter blueprint crosses this line.

1.5 Methodology (ten steps)

Observe → Transcribe → Extract → Structure → Agentise → Test → Shadow → Controlled Pilot → Measure → Expand responsibility.

Starter gets you through Structure → Shadow with templates. Professional / Team deepen Test, Measure, Expand.

2. Human-vs-agent decision framework

Use before writing any instruction set.

2.1 Decision grid

For each task class, mark **H** (human only), **A-assist** (agent drafts), **A-bounded** (later L2+), or **Out**.

Is the task frequent and rules-partially-known?	Keep human / improve process first
Can success be checked against evidence (PO, GR, policy)?	Do not agentise judgement-only politics
Is failure recoverable without payment leakage?	Keep L0–L1 or Out
Is SoD preserved if agent prepares?	Redesign
Can we log inputs, rule, confidence, human decision?	Do not deploy
Does policy already define tolerances / DOA?	Write policy first

2.2 Default placements

Capture & classify inbound invoice	A-assist
Validate completeness fields	A-assist
Suggest match / explain variance	A-assist
Force-match above tolerance	H
Classify exception code	A-assist → A-bounded later
Clear material exception	H (agent prepares pack)

Draft supplier / requester chase	A-assist (human send)
Approve invoice under DOA	H (agent routes only)
Flag duplicate candidates	A-assist
Confirm duplicate / fraud	H
Prepare payment proposal	A-assist
Authorise / release payment	H always
Approve vendor bank change	H always
Period-end assertion sign-off	H always

2.3 “Agentise later” list

Park tasks that fail the grid. Improving PO quality or GR discipline often beats a clever agent on a broken upstream.

3. Top 10 agent blueprints (condensed)

Each blueprint is usable. Expand fields in Professional charters when you scale.

Shared blueprint fields

- Job · Owner role · Inputs · Outputs · Does · Does not · Escalation · Starter autonomy · KPIs · Evidence

A01 — Invoice Intake

Job	Capture, classify, and register inbound invoices into a controlled work queue.
Owner	AP Intake Lead
Inputs	Email/portal/EDI payloads; channel rules; vendor directory
Outputs	Registered invoice stub; channel tag; confidence; rejection reason if any
Does	Detect invoice vs non-invoice; extract header candidates; route by entity/channel; log duplicates *candidates* to A10
Does not	Drop invoices silently; invent vendor masters; approve payment
Escalation	Unreadable / wrong entity / malware suspicion → human intake + IT security path
Starter autonomy	L0–L1
KPIs	Time-to-register; % auto-classified correctly (sampled); silent-drop incidents (=0)
Evidence	Source message ID; extract JSON/fields; human correction log

A02 — Invoice Validation

Job	Verify completeness and coding readiness before match.
Owner	AP Validation Lead
Inputs	Invoice image/fields; tax rules summary; vendor master; PO-required policy
Outputs	Pass / fail with coded gaps; validation checklist
Does	Check supplier identity signals, invoice #, dates, amounts, tax fields, PO presence rules, attachment presence
Does not	Invent commercial terms; override PO policy; post GL
Escalation	Identity conflict / banking detail on invoice → enhanced human path
Starter autonomy	L0–L1
KPIs	First-pass validation rate; false-pass rate (QA sample); recycle loop count
Evidence	Checklist; failed rules; reviewer decision

A03 — Matching

Job	Perform PO / non-PO match logic; explain variances.
Owner	AP Match Lead
Inputs	Invoice lines; PO; GR; tolerances
Outputs	Match status; variance pack; suggested next agent
Does	Two-/three-way compare within tolerance; explain price/qty gaps; route exceptions to A04
Does not	Force-match beyond tolerance; change PO price; waive DOA
Escalation	Material variance / chronic vendor → Buyer + Category
Starter autonomy	L0–L1 (L2 only with Pro controls)
KPIs	Match cycle time; within-tolerance auto-suggest accept rate; force-match attempts (=0 by agent)
Evidence	Line comparisons; tolerance rule ID; human disposition

A04 — Exception Triage

Job	Classify exceptions, assign owners, start SLA clocks, route resolvers.
Owner	Exception Desk Lead / Head of AP
Inputs	Failed validation/match signals; taxonomy; ownership matrix
Outputs	Primary code; owner; due; linked resolver agent
Does	Apply taxonomy; detect multi-code; prioritise by age×value×risk; prevent unowned WIP
Does not	Clear exceptions; invent codes; hide aged items
Escalation	Unowned >48h material → Process Owner
Starter autonomy	L0–L1
KPIs	% coded within 24h; unowned count; reclassify rate
Evidence	Code; rationale; timestamp; owner ack

A05 — Goods Receipt

Job	Chase and reconcile receipt evidence for receipt-required invoices.
Owner	GR Coordination (AP) with Ops receivers
Inputs	PO; invoice qty; GR search; receiver list; delivery notes if any
Outputs	Chase pack; GR found/missing status; partial receipt flag
Does	Locate likely receiver; draft chase; reconcile partials; update A04
Does not	Post GR without policy; fake receipt dates
Escalation	Supply-critical / due date risk → Ops supervisor
Starter autonomy	L1 drafts; human send
KPIs	GR chase response time; % resolved by GR post; false chase rate
Evidence	Chase thread; GR doc #; human confirmation

A07 — Approval (routing assist)

Job	Drive approval routing within DOA; never invent authority.
Owner	AP Workflow Lead
Inputs	Invoice; DOA table; approver directory; coding

Outputs	Correct approver queue; nudge drafts; stale approval alerts
Does	Map amount/entity to DOA; detect missing approver; remind with evidence pack
Does not	Approve; reassign above policy; bypass SoD
Escalation	DOA gap / conflict → Controllership
Starter autonomy	L1
KPIs	Approval cycle time; misroute rate; bypass attempts (=0)
Evidence	DOA rule applied; approver ID; timestamps

A08 — Supplier Resolution

Job	Coordinate supplier queries with human send/approve gates.
Owner	Supplier Query Lead
Inputs	Exception pack; vendor contacts; prior correspondence
Outputs	Draft email/portal message; response log; proposed resolution
Does	Clarify price/qty/credit/duplicate questions; attach evidence; track promises
Does not	Commit payment dates that override policy; accept bank changes from email alone
Escalation	Dispute / legal tone → Procurement + AP Lead
Starter autonomy	L1 (human send)
KPIs	Response SLA; reopen rate; tone QA fails
Evidence	Message + attachments + human approval to send

A09 — Internal Follow-Up

Job	Chase requesters, buyers, GR clerks with evidence packs.
Owner	AP Exception Coordinator
Inputs	Exception code; named internal owner; SLA; documents
Outputs	Chase sequence; escalation ladder; resolution notes
Does	Timed nudges; escalate by role; stop when answered

Does not	Spam; escalate personally; clear without evidence
Escalation	Per matrix (e.g. 48h → manager)
Starter autonomy	L1
KPIs	Internal response time; escalations hit rate; courtesy QA
Evidence	Pack + timestamps + human overrides

A10 — Duplicate & Anomaly

Job	Surface duplicate and anomaly candidates ; prepare investigation packs.
Owner	AP Controls Lead
Inputs	Invoice register; payment history; similarity features
Outputs	Candidate pairs; score explanation; recommended hold
Does	Exact and near-match detection; hold suggestion; pack for human
Does not	Guarantee fraud/duplicate detection; auto-confirm without policy; release payment
Escalation	Critical score + material \$ → Controls same day
Starter autonomy	L0–L1
KPIs	Precision/recall on samples (local); time-to-disposition; missed confirmed duplicate after clear (incident)
Evidence	Feature list; decision; payment status at decision

A11 — Vendor Statement Reconciliation

Job	Reconcile supplier statements to open items; propose discrepancy packs.
Owner	AP Reconciliation Lead
Inputs	Statement; AP open items; remittances
Outputs	Matched lines; discrepancies coded; draft supplier queries
Does	Line match; detect missing invoices / missing payments / timing
Does not	Auto-settle disputes; book adjustments without human
Escalation	Large unreconciled → AP Lead + Procurement
Starter autonomy	L0–L1

KPIs	Statement cycle time; unreconciled \$; pack accept rate
Evidence	Recon worksheet; codes; sign-off

A12 — Payment Proposal Review

Job	Prepare payment batches for human authorisation only.
Owner	AP Payments Lead (prep); Treasury/AP leadership (auth)
Inputs	Approved open items; holds; duplicate flags; early-pay rules
Outputs	Proposal list; exception flags; exclusion list
Does	Assemble candidates; strip holds; highlight anomalies for review
Does not	Authorise; generate/transmit bank file as autonomous act; override holds
Escalation	Ambiguous inclusion → human; never retry ambiguous sends
Starter autonomy	L1 prepare; auth H
KPIs	Proposal prep time; items incorrectly included (QA); auth bypass attempts (=0)
Evidence	Proposal version; flags; authoriser ID; timestamp

4. Exception taxonomy (condensed — all major codes)

Use **one primary code** per case. Full attribute sheets live in Professional; Starter gives operating definitions.

`missing_po`	Required PO absent/blank	Requester / Buyer	A02, A04, A09, A06
`invalid_po`	PO ref invalid / wrong vendor-entity	Buyer	A03, A06, A09
`po_closed`	PO closed/cancelled while invoice live	Buyer	A06, A09
`po_exhausted`	Remaining PO insufficient	Buyer	A06, A03, A10
`price_mismatch`	Price beyond tolerance vs PO/contract	Buyer / Category	A03, A06, A08
`quantity_mismatch`	Qty beyond tolerance vs PO/GR	Receiver / Buyer	A03, A05, A08
`missing_receipt`	3-way required; no GR	Receiver	A05, A09

`partial_receipt`	GR covers only part of invoice	Receiver / Supplier	A05, A03, A08
`duplicate_invoice`	Confirmed duplicate	AP Controls	A10, A04, A12
`potential_duplicate`	Candidate not yet confirmed	AP Processor / Controls	A10, A12
`wrong_supplier`	Supplier identity / remit mismatch	AP + Master Data	A02, A08, A10
`incorrect_legal_entity`	Wrong company code / entity	AP / Requester	A02, A09
`tax_issue`	Tax fields / treatment fail checks	Tax / AP	A02
`approval_missing`	Required approval not complete	Approver	A07, A09
`doa_issue`	DOA conflict / insufficient authority	Controllershship	A07
`coding_missing`	Cost object / GL incomplete	Requester	A02, A09
`invalid_cost_centre`	Cost object invalid / closed	Finance / Requester	A02, A09
`invoice_quality`	Illegible / incomplete document	Supplier / Intake	A01, A08
`ocr_extraction_issue`	Capture fields wrong	AP Intake / Systems	A01, A02
`master_data_issue`	Vendor master blocking	Master Data Steward	A02, A08
`banking_change_concern`	Remit/bank change signal	Controls / Treasury	A10, A08 — human verify
`credit_note_required`	Credit expected to resolve	Supplier / Buyer	A08, A03
`statement_discrepancy`	Statement vs subledger gap	AP Rec	A11
`payment_hold`	Policy/control hold applied	AP Controls	A12, A04
`disputed_invoice`	Commercial/legal dispute	Procurement / Legal	A08 — human lead
`aged_unresolved_item`	Past SLA; still open	Process Owner	A04, A16
`system_interface_error`	Integration/posting technical fail	IT / AP Systems	A01, A16

Operating rules

- Prefer specificity over `aged\_unresolved\_item` (use as secondary).
- `potential\_duplicate` stays until human disposition.
- `banking\_change\_concern` never auto-clears.
- Tune SLAs locally; do not rename codes casually.

5. AI readiness diagnostic summary

Use Free Diagnostic for full 32 questions. Starter summary:

Emerging (0–39)	Foundations weak	Taxonomy + controls first; agents L0 only
Structured (40–59)	Operable	One L0→L1 pilot + baselines
Agent-Ready (60–79)	Ready to design	≤3 agents; L1 ceiling; weekly scorecard
Operating (80–100)	Evidence culture	Expand with gates; still no autonomous pay

Minimum before L1 recommend on live work: coded exceptions (or pilot slice), named owner, exclusions written, baseline attempt, human payment auth verified.

6. Process-mapping template (how to fill)

See also 02_TEMPLATES_PACK.md.

- 1 Pick **one** slice: entity × channel × category (e.g. US email × inventory PO).
- 2 List steps from inbound to payment proposal.
- 3 Mark each step: Human / System / Candidate agent.
- 4 Note artefacts (invoice, PO, GR, approval).
- 5 Note failure modes → taxonomy codes.
- 6 Mark control points (SoD, holds, dual auth).
- 7 Identify data fields required for agent inputs.
- 8 Stop at payment **proposal**; draw a hard line before authorisation.

7. Agent job description template (guidance)

A job description is the **employment contract** for the agent — not a prompt dump.

Must include: purpose; scope in/out; responsibilities; exclusions; human owner; RACI; autonomy level; escalation; success metrics; review cadence.

Fill blanks in Templates Pack. Approve with Head of AP before any L1.

8. Agent instruction template (guidance)

Instructions operationalise the job:

- 1 Objective (one paragraph)
- 2 Tools / systems allowed
- 3 Input contract (required fields)
- 4 Decision procedure (ordered checks)
- 5 Output schema
- 6 Confidence / abstain rules
- 7 Escalation triggers
- 8 Tone rules (if communicating)
- 9 Prohibitions (hard)
- 10 Logging requirements

Abstain > guess on identity, banking, duplicates, and DOA.

9. KPI scorecard (starter)

Track weekly for each pilot agent + AP slice.

Volume touched						
% coded exceptions						
Median cycle time						
Human accept rate of recommendations						
Critical control incidents						0
Silent drops / unlogged actions						0
QA sample pass %						Local threshold
Operator trust (1–5 survey)						

Demote or pause if critical control incidents >0 or accept rate collapses without process cause.

10. Governance checklist (Starter)

1	Executive sponsor named	CFO / Controller	■
2	AP Process Owner named		■
3	Agent Owner named per pilot agent		■
4	Autonomy documented (≤L1 recommended) ceiling		■
5	Payment auth remains human — confirmed in writing		■

6	Vendor bank change path excluded from agent approve		■
7	Exception taxonomy adopted for pilot slice		■
8	SoD matrix reviewed for agent prepare roles		■
9	Logging / retention path defined	IT / AP Systems	■
10	Shadow plan + QA sample size set		■
11	Kill-switch / pause owner named		■
12	Change control for instructions / prompts		■
13	No confidential data in public AI tools		■
14	Internal Audit informed (awareness)		■
15	Rollback plan (<1 hour demotion)		■

11. AP transformation roadmap — 90-day sketch

0 Foundation	1–15	Diagnostic, taxonomy, baselines, scope slice	Codes live; baseline sheet; ceiling agreed
1 Design	16–35	Maps, job descriptions, instructions, RACI	Charters approved; checklist green
2 Shadow	36–60	L0/L1 shadow on slice; daily exception huddle	Shadow agreement sampled; zero critical incidents
3 Controlled pilot	61–75	L1 recommendations in operating rhythm	Accept rate stable; evidence packs retained
4 Decide	76–90	Steering: hold / expand agents / buy Professional depth	Written decision + next quarter backlog

Not in 90 days: fleet-wide L3, autonomous payments, ERP replacement, guaranteed FTE reduction claims.

12. Implementation checklist

Pre-pilot

- ☐ Scope one-pager approved
- ☐ Taxonomy in tracker

- ☐ Baselines captured or explicitly “unknown → finding”
- ☐ Charters v1 approved
- ☐ Access read-only proven
- ☐ QA samplers trained
- ☐ Comms to AP team (what changes / what does not)

During shadow

- ☐ Daily volume log
- ☐ Disagreement log (agent vs expert)
- ☐ Incident log (even near-misses)
- ☐ Weekly 30-min review

Before any L2 discussion

- ☐ Professional control matrix reviewed
- ☐ Controller co-sponsor
- ☐ ≥4 weeks L1 evidence
- ☐ Payment boundary re-confirmed

13. When to upgrade

Full 16-agent charters, governance framework, control matrix, business-case model, UAT	Professional — US\$199
Workshop, training, executive packs, trackers, multi-user facilitation	Team — US\$499
Facilitated blueprint for your stack	Custom Blueprint — US\$1,500–3,000

14. Licence note (individual)

Grant: One named individual may use this Starter Kit for professional purposes inside **one** employer / client organisation at a time.

Includes: Personal use of templates; internal workshops the licensee facilitates alone.

Does not include: Sharing files with a whole shared-services centre as a team standard; resale; public redistribution; white-labelling; uploading the kit as the system prompt library for a commercial product.

For multi-user / workshop licence rights, purchase **Team**.

15. Disclaimers

No guaranteed savings, fraud detection, compliance certification, accounting accuracy, or ROI. Not legal or audit advice. Payment authorisation stays human. Local policy overrides templates.

Evidence Room — Operating Evidence. *ER-START-GUIDE-001 v1.0 · Individual licence*