

AP Agent Starter Kit

The \$79 kit that should feel like a \$500 briefing

Product	AP Agent OS
Tier	Starter
Version	1.0
Date	September 2026
Use	Implementation toolkit

Responsibility is earned.

Start here

This kit is enough to redesign one AP agent slice without buying the full OS.

Use it in this order:

- 1. Operating model (this folder)
- 2. Human vs agent decision
- 3. Top 10 agent blueprints
- 4. Exception taxonomy (starter)
- 5. Process map template
- 6. Agent job description + instruction templates
- 7. KPI scorecard (lite)
- 8. Governance checklist
- 9. 90-day roadmap and implementation checklist

If you need all 16 agents, control matrix, ROI model, UAT, shadow methodology and executive packs, upgrade to Professional.

Operating model in one page

AP work has four layers:

Layer	Human default	Agent default
Policy and DOA	Own	Never set
Payment authorisation	Own	Review only, never release
Exception judgment	Own	Recommend / prepare
Repeatable analysis, drafting, routing, monitoring	Review	Observe → recommend → prepare

Agents are roles, not features. Each role has a charter, owner, exclusions, KPIs, and an autonomy level that starts at 0 or 1.

Human vs agent decision

Score a task 0–2 on each: frequency, rules clarity, data completeness, reversal cost, regulatory sensitivity, supplier impact.

- Deterministic automation if the rule is stable and the data is structured (classic match within tolerance).
- Agent recommend if language, judgment, or incomplete data is involved.
- Human only if the action posts, pays, changes bank details, or grants access.

What can go wrong

Do not use an LLM to decide a three-way match when the ERP already can. Agents should not re-implement match engines.

Ninety-day starter roadmap

Days	Outcome
1–15	Diagnostic, baseline KPIs, one process map
16–30	First charter, controls, historical test pack
31–60	Shadow mode, weekly scorecard
61–90	Controlled recommend-mode pilot or a clear stop

Actual duration depends on systems, data quality, integrations and governance. The 4–6 week “one agent” story is only for a well-bounded, data-ready use case.

This material is an implementation toolkit. It is not legal, tax, accounting, audit or investment advice. Agents do not authorise payments. Humans remain accountable. Statistics are cited in the Research Ledger.