

Testing, Shadow Mode and Pilot

Product	AP Agent OS
Tier	Professional
Version	1.0
Date	September 2026
Use	Implementation toolkit

Responsibility is earned.

Historical testing

1. Draw 50–100 cases from the last 90 days for the first agent. Stratify: easy, typical, nasty, high-value, high-risk codes.
2. Create a gold outcome (human dual-review on 20% if you can).
3. Freeze model, prompt, taxonomy versions.
4. Run. Score with the KPI dictionary.
5. Read every false negative on high-risk codes.
6. Do not proceed to shadow if you cannot explain the misses.

Pass guideline (first agent, Level 1): classification or recommendation accept rate agreed in advance (example 85% on the typical band), and no unexplained high-risk false negative. These numbers are local gates, not industry laws.

UAT

UAT is not a demo. Scripts in the Excel/Doc templates cover: happy path, known exception, hostile invoice text, missing data, model timeout, override, fallback.

Sign-off: agent owner, Controls, tool owner. Head of AP if any Level 2+ send/write.

Shadow mode

- Agent writes to a parallel queue only.
- Humans ignore it for processing; a reviewer compares daily.
- No supplier email, no ticket in the real workflow unless labelled test.
- Minimum two weeks or 200 live items, whichever is later for high-volume teams.
- Exit: written compare, updated charter, go/no-go.

Controlled execution / pilot

Limit entity, category, value, named users, calendar. Rollback: disable send/write in one step. Communications: tell the affected buyers/receivers what will look different.

Promotion

Use the governance evidence pack. Expiry 90 days.

Implementation phases (OS view)

Phase	Name	Typical artefact
0	Baseline and readiness	Diagnostic, KPI baseline
1	Process discovery	Observation + map

Phase	Name	Typical artefact
2	Agent specification	Charter
3	Data / tool access	Access list, SoD
4	Prototype	Sandbox outputs
5	Historical testing	Scored pack
6	Shadow	Shadow log
7	Controlled execution	Pilot charter
8	Performance review	Weekly report
9	Responsibility progression	Autonomy record
10	Scale	Next agent or next entity

4–6 weeks is an illustrative path for one well-bounded Level 1 agent with data already extractable. It is not a commitment.

This material is an implementation toolkit. It is not legal, tax, accounting, audit or investment advice. Agents do not authorise payments. Humans remain accountable. Statistics are cited in the Research Ledger.