

TIER 3 · US\$499 · FIVE NAMED SEATS · ONE LEGAL ENTITY

AP Agent OS — Team Edition

A transformation leader can run a session from this pack tomorrow.

Evidence Room — AP Agent OS Team Edition

00 — Team Edition Playbook

Product: Evidence Room AP Agent OS — Team Edition

Brand idea: Proof before permission. Agents earn responsibility. Evidence decides.

Audience: Finance Transformation lead, Head of AP / Shared Services, Controller, Change lead, Control owner

ERP stance: Agnostic.

Examples: ACME Manufacturing — **ILLUSTRATIVE** only. Buyer fields [BUYER] .

Version: 1.0

Classification: Enterprise transformation pack. Complements Professional OS. Does **not** promise ROI, fraud reduction, compliance, or autonomous payments.

What Team Edition is

Professional OS is the **rulebook**. Team Edition is the **way a group of humans adopts the rulebook**: workshop, interviews, training, steering, benefits tracking, maturity, and change communications.

You can run a one-day workshop from this pack tomorrow **if** you have: a named AP Accountable, a room (or video) for 8–16 people, exported volume/exception samples (even imperfect), and a first-slice hypothesis. You cannot leave the day with a live executing agent — and you should not try.

Use Professional files as the authority when Team materials summarise them.

Need	Start here
Run a day	Workshop/02_WORKSHOP_AGENDA_ONE_DAY.md + 01_FACILITATION_PACK.md
Exercises and outputs	Workshop/03_WORKSHOP_EXERCISES.md
Before the day	Workshop/04_STAKEHOLDER_INTERVIEW_GUIDE.md ; 05_PROCESS_OWNER_QUESTIONNAIRES.md
Train processors	Training/01_AP_TEAM_TRAINING.md
Train agent owners	Training/02_AGENT_OWNER_TRAINING.md
Steer the programme	Executive/01_STEERING_COMMITTEE.md
Brief the CFO	Executive/02_CFO_BRIEF.md
Brief the board	Executive/03_BOARD_SUMMARY.md
Track work	Implementation/01_IMPLEMENTATION_TRACKER.md
Track benefits honestly	Implementation/02_BENEFITS_REALISATION.md

Need	Start here
Score maturity	Implementation/03_MATURITY_ASSESSMENT.md
Change the organisation	Change_Management/01_CHANGE_TOOLKIT.md ; 02_EXECUTIVE_COMMS.md
Operating rules	../03_AP_AGENT_OS_PRO/00_PROFESSIONAL_OPERATING_SYSTEM.md

Every section below: **what / how / who / wrong / control / measure / evidence.**

1. What to do with this pack

1. Appoint a **Transformation lead** (Responsible for the pack) and a **Head of AP** (Accountable for AP outcomes).
2. Run pre-workshop interviews (2–5 days).
3. Run the one-day workshop; leave with a signed first-slice and a 30-day evidence plan — not a slogan.
4. Train the floor before Shadow.
5. Install steering and a benefits register that cannot accept unsigned cash.
6. Re-assess maturity each quarter.
7. Promote agents only through Professional gates.

How. Sequence in §3.

Who.

Role	Team Edition duty
CFO / FD	Sponsor; L3+ / payment-adjacent acknowledgement
Controller	Financial claims; cost model
Head of AP	Operating Accountable
Transformation lead	Workshop, tracker, comms cadence
AP Manager	Daily evidence after Test
Control owner	Veto on risk-control
Change lead	Toolkit and floor messaging
Finance Systems	Identities and hashes
Internal Audit	Informed; may observe workshops
Agent	Never Accountable

Wrong. “Centre of Excellence” owns the agent; AP is a passenger. Vendor success manager as owner.

Control. RACI from Professional. Incumbents on the tracker.

Measure. Workshop outputs filed; training completion before Shadow; steering packs with dictionary IDs.

Evidence. [BUYER]/Evidence/Team/ plus Professional folders.

2. Outcomes Team Edition is allowed to claim

Allowed: We named owners. We froze a slice. We produced a current-state gap list. We trained N people. We have a Test pack ID. We held Shadow for N days. We decided Hold.

Not allowed: Guaranteed savings. "Autonomous AP." "Fraud-proof." "Compliant by design." "Board-ready ROI" without Controller validation.

Context research (not our results): Deloitte 8 Oct 2025 — 63% AI deployed / 21% measurable ROI; Hackett Nov 2025 — software-adopter 60% touchless / 3.5× productivity at ≥30% touchless. Use only as labelled context in executive papers.

3. How — 90-day team path (ILLUSTRATIVE)

Aligns with Professional Phase 0–7 for one slice; Pilot may start after day 90.

Window	Team activity	Professional phase
Days 1–10	Interviews, questionnaires, mobilise, cost-model start	0–1
Day 11	One-day workshop	0–4 drafts
Days 12–25	Structure + charter + risk + baseline	2–4, 7 start
Days 18–40	Environment (often the slip)	5
Days 26–45	Historical test + UAT	6, UAT
Days 40–90	Shadow + training reinforcement + first steer	7
After evidence	Pilot conversation	8

If IdP/ERP access slips, **do not** compress Test.

4. Workshop-in-one-page (detail in Workshop/)

Object of the day: a first-slice pack the Accountable will sign.

Not the object: a 16-agent roadmap poster; a savings target; a tool selection bake-off (tools are a later C2 with privacy).

Must-produce artefacts by 17:00

1. Slice allow-list
2. Human-class confirmation

- 3. As-done vs as-written top 10
- 4. First agent + L0 charter skeleton
- 5. Risk top 8
- 6. Baseline metric list (IDs, not targets)
- 7. 30-day plan with named R's
- 8. Decision: proceed to Observe completion / Hold / Stop

Who in the room (minimum)

Head of AP, AP Manager, one processor, one exception owner, Control, FinSys, Transformation. Add Procurement and GR if matching/GR is the candidate. Add Treasury if anyone mentions payments.

5. Training-in-one-page

Processors learn: packet fields, taxonomy, override reason codes, "do not wait in Shadow," kill-switch, four human classes.

Owners learn: charter checklist, promotion evidence, weekly report, envelope, incident.

Neither course is a prompt-engineering class.

6. Steering-in-one-page

Monthly (or per gate). Pack order: risk-control → operational → cost → activity last. Conservative business case if still funding. Kill-switch test date. No vendor theatre.

7. What can go wrong (programme)

Failure	Control
Workshop becomes a vision offsite	Artefact list; Decision page
Change team writes "AI will free you for analysis" as a promise	Toolkit language list
Benefits register copies Upside	Benefits file rules
Maturity scored by enthusiasm	Evidence column required
Team Edition used to bypass Test	Playbook §1 step 7

8. How you measure Team Edition itself

Question	Metric
Can they run Monday?	First-slice signed; owners named
Did the floor hear the rules?	Training attendance before Shadow
Is Steer honest?	Unsigned \$ lines = 0
Did we over-claim?	Language review of comms

Evidence. Team scorecard quarterly (qualitative + the metrics above).

9. ACME Manufacturing — ILLUSTRATIVE programme

Fictional ACME: 8,400 invoices/month, shared AP, ERP + mailbox. Workshop chose Intake + Duplicate, ACME US, email PDF, L0. Did not choose Payment Review as first agent. 90-day path reached Shadow; Pilot refused until paid-sample FNR reviewed. No board ROI slide.

10. What this playbook is not

A substitute for Professional OS. A licence to implement on a non-approved tool. A certificate.

Proof before permission. Agents earn responsibility. Evidence decides.

End of 00_TEAM_EDITION_PLAYBOOK.md

Evidence Room — AP Agent OS Team Edition

Workshop — 01 Facilitation pack

Product: Evidence Room AP Agent OS — Team Edition

Module: Workshop

Standard: Proof before permission

Audience: Transformation / Process Excellence facilitator

Version: 1.0

This pack lets you facilitate tomorrow. Pair with `02_WORKSHOP_AGENDA_ONE_DAY.md` and `03_WORKSHOP_EXERCISES.md`.

1. What to do

Design the room, pre-work, materials, roles, and parking-lot rules so the day produces **signed artefacts**, not a mural.

How. Checklists below.

Who. Facilitator Responsible; Head of AP Accountable for decisions taken in the room; Control owner may halt a path that implies execute-on-payments.

Wrong. Facilitator becomes the process owner. Vendor runs the day.

Control. Decision page; artefact checklist at 16:30.

Measure. Artefacts complete / 8; decision enum recorded.

Evidence. [BUYER]/Evidence/Team/Workshop/[DATE]/ .

2. How — 10 days before

Task	Owner	Done
Confirm Head of AP attends the decision blocks (not only kick-off)	Transformation	☐
Invite list (see §4)	Transformation	☐
Send questionnaires (05_PROCESS_OWNER_QUESTIONNAIRES.md)	Transformation	☐
Book 2–4 stakeholder interviews (04_...)	Transformation	☐
Request extracts: volume by channel/type (90 days), top exception reasons, open GR/IR age, payment calendar	AP Manager	☐
Confirm recording / photo policy	Privacy / AP	☐
Room: wall space or digital board with export	Facilitation	☐
Print or share Professional one-pagers: human classes, packet fields, L0–L4	Facilitation	☐
Draft first-slice hypothesis (to attack, not to rubber-stamp)	Head of AP + Transformation	☐
Approved-tool reminder: no live personal chatbot in the room	FinSys	☐

If extracts fail, still run the day on **interviews + operator testimony**, and mark volume cells ASSUMED.

3. How — materials

Must have

- Agenda (02) on one page

- Exercise cards (03)
- Blank slice sheet, RACI translation, risk catalogue (.. /03_AP_AGENT_OS_PRO/Templates/)
- Exception taxonomy index
- Sticky notes or digital equivalent in **four colours**: Fact / Assumption / Decision / Parking
- Decision page (Templates/MEETING_GUIDE.md)
- Kill-switch and four-human-classes poster

Nice to have

- 10 redacted invoice packets (real buyer data, privacy-cleared)
- ACME ILLUSTRATIVE pack only if buyer data is forbidden — label it

Do not bring

- A savings target
- A 16-agent implementation Gantt for sign-off
- Vendor pricing as the frame of the day

4. Who — room design

Seat	Required?	If absent
Facilitator	Yes	Do not run
Head of AP	Yes for 09:00, 12:30, 16:00	Decisions are provisional
AP Manager	Yes	
Processor (as-done witness)	Yes	Discovery is fiction
Exception owner	Strongly	
Control owner	Yes	No charter skeleton
Finance Systems	Yes	Phase 5 fantasy
Transformation	Yes	
Controller or FP&A	Half-day or 15:00	Cost/baseline weak
Procurement / GR	If those agents are candidates	
Treasury	If payment mentioned	Parking lot becomes a veto later
Internal Audit	Optional observer	
Vendor	No (or last 20 min demo only)	Avoid capture of the agenda

Maximum 16. Split if more.

SoD note: payment authoriser should not sit as “the AP redesign lead.”

5. How — facilitator behaviours

Do

- Ask “what evidence would we file?” after every proposal.
- Separate Fact / Assumption with colours.
- Read the human-classes line aloud when someone says “just pay it.”
- Timebox; park strategy debates.
- Quote operators verbatim onto the as-done list.

Do not

- Rescue a bad slice because the calendar needs a win.
- Let the senior person overwrite the processor without logging dissent.
- Translate “we usually tolerate a couple of percent” into a number.
- Allow “the AI will handle exceptions.”
- Promise ROI.

Interventions (scripts)

You hear	You say
“Let’s automate all of AP this quarter.”	“Professional START HERE is one increment. Write the allow-list or we stop at parking.”
“We’ll skip testing — it’s low risk.”	“Then it stays L0. Test is the permission mechanism.”
“Hackett says 60% touchless.”	“That is software-adopter research, not our baseline. Our STP is [extract or unknown].”
“The agent can approve if confidence is high.”	“Confidence cannot widen DOA. That’s a control break.”
“We’ll use my ChatGPT login.”	“Out of order. Approved-tool list or we do not run live work.”

6. What can go wrong

Failure	Control
Happy-path theatre	Processor seat; redacted ugly invoices
Decision by the loudest	Decision-maker named on the page
Parking lot as a second workshop	Lot has owners and dates before close
Digital board not exported	Facilitator exports before leaving the call
Informal promises in the corridor	“If it is not on the decision page it did not happen.”

7. Measure the day

Metric	Target discipline
Artefacts 1–8 complete	Count; missing = Hold
Assumptions labelled	Share of slice cells
Processor talk time	Facilitator notes; if ~0, day is invalid
Decisions with Accountable	100% of decision page rows

Evidence. Board export, photos, decision page, attendance, questionnaires received.

8. Day-of run sheet (facilitator only)

Clock	You are doing
–30 min	Boards named; colours; packs at seats
08:55	Headcount; who is missing
Each exercise	Read the object aloud; set timer; do not join the content
Every break	Photograph walls
16:20	Artefact completeness check — name the holes
16:40	Decision page — Head of AP writes, not you
+30 min	Export, file, send recap (facts and decisions only)

9. Recap email (send same day)

Subject: AP Agent OS workshop [DATE] — decisions and evidence

Body (adapt): slice; first agent; autonomy **L0**; out-of-scope verbs; 30-day R's; decision enum; folder path; **no savings sentence**.

Proof before permission.

End of 01_FACILITATION_PACK.md

Evidence Room — AP Agent OS Team Edition

Workshop — 02 One-day agenda

Product: Evidence Room AP Agent OS — Team Edition
Module: Workshop
Standard: Proof before permission
Audience: Participants and facilitator
Version: 1.0
Duration: 08:30–17:00 local (ILLUSTRATIVE). Compress only by dropping the optional 14:30 cost block — **do not** drop human classes or decision.

Times include movement. Exercises are specified in 03_WORKSHOP_EXERCISES.md (E1–E8).

1. What to do

Run the day as a **working session**. Phones down for exercise blocks. Outputs are written artefacts.

How. This timetable. Facilitator uses 01_FACILITATION_PACK.md .

Who. See facilitation pack §4. Head of AP must be present at blocks marked **DECIDE**.

Wrong. Morning keynote + afternoon “working groups” with no plenary decision.

Control. Timekeeper; artefact list on the wall all day.

Measure. Start on time; exercises finished or explicitly cut with Accountable consent.

Evidence. Timed agenda initialled; board export.

2. Objective (read at 08:30)

By 17:00 we will have a **first-slice allow-list**, a **first agent at L0**, a **risk and baseline list**, and a **30-day evidence plan** — or a documented **Hold/Stop**. We will not select a vendor, set an ROI target, or grant execution rights.

3. Timetable

Time	Block	Exercise	Object	Who must be in the room
08:30–08:50	Open	—	Rules, human classes, decision rights	All

Time	Block	Exercise	Object	Who must be in the room
08:50–09:20	Context	E1	Facts vs assumptions on volume and pain	All; AP Manager brings extracts
09:20–10:20	As-done	E2	Top unofficial rules and money-adjacent map	Processor + exception owner speak first
10:20–10:35	Break		Photograph walls	
10:35–11:20	Slice	E3	Allow-list + exclusions	Head of AP DECIDE
11:20–12:10	Agent	E4	First agent, L0–L4 sanity, forbidden tools	Control + AP Manager
12:10–12:40	Human line	E5	Four classes + SoD pairs on this slice	Control DECIDE halt if payment execute appears
12:40–13:20	Lunch		Do not redesign at the sandwich table without notes	
13:20–14:10	Risk	E6	Top 8 risks, owners, “Pilot blocker?”	Control
14:10–14:50	Measure	E7	Baseline KPI IDs; no targets from blogs	AP Manager + FP&A if present
14:50–15:05	Break			
15:05–16:00	30-day plan	E8	Named R’s, Phase 0–6 tasks, caveats	FinSys for access tasks
16:00–16:25	Playback	—	Facilitator reads artefacts 1–8; holes named	All
16:25–16:50	DECIDE	—	Decision page: Proceed / Hold / Stop	Head of AP writes
16:50–17:00	Close	—	Recap path; training date; store path	All

Optional if Controller attends 14:10: five minutes on Conservative vs Base vs Upside **as a method**, not as numbers to applaud.

4. Opening script (08:30, 3 minutes — facilitator)

This is Evidence Room Team Edition. Proof before permission. Agents earn responsibility. Evidence decides.

We will not promise savings. We will not call this fraud detection or a compliance programme.

Payment release, vendor bank changes, policy exceptions, and legal disputes stay human.

The Head of AP is Accountable for decisions today. I am not.
If you state a number, say whether it is extracted, estimated, or illustrative.

5. Parking-lot rules

Parking items get an **owner and a date** at 16:00 or they are deleted.

Forbidden parking: “define fraud approach,” “pick the LLM,” “set 30% efficiency.” Those are either out of product scope or belong in Test/Measure, not a lot.

6. Remote / hybrid variant

Same times. Two facilitators if >8 on video (content + board). Processor must have a working microphone. Breakout rooms only for E2/E8, 15 minutes max, then plenary. Export the board before anyone leaves.

7. Half-day emergency cut (not preferred)

If you only have 4 hours: E1 (15), E2 (40), E3 (35), E5 (25), E8 (40), DECIDE (20). You **must** schedule E4/E6/E7 within five working days or the slice is not Agentise-ready.

Wrong. Declaring the programme launched after a half day.

Control. Calendar hold for the missing exercises before people leave.

8. What can go wrong

Failure	Control
Sponsor speech eats 08:30–09:30	Sponsor gets 5 minutes inside Open
E2 skipped because “we know our process”	Processor seat mandatory
Decision rushed at 16:55	Playback starts 16:00 even if E8 is ugly

9. Measure / evidence

Attendance; start/end actuals; eight artefacts; decision page; photos.

Proof before permission.

End of 02_WORKSHOP_AGENDA_ONE_DAY.md

Evidence Room — AP Agent OS Team Edition

Executive — 02 CFO brief

Product: Evidence Room AP Agent OS — Team Edition

Module: Executive

Standard: Proof before permission

Audience: CFO / Finance Director (and Controller as preparer)

Version: 1.0

A 10-minute read. Replace [BUYER] and any ILLUSTRATIVE ACME lines with buyer facts. Do not insert guaranteed ROI.

1. What to do

Send this brief before first Steer or funding. The CFO should leave knowing: what we are buying (a **permission system** for AP agents), what we will not do (autonomous payments), how money will be talked about, and what decision is asked.

How. Sections 2–8 as the memo body.

Who. Controller or Head of AP authors; Transformation may draft; CFO is the reader.

Wrong. A vision deck. A vendor TCO as the brief.

Control. Controller reviews any number. Research labelled.

Measure. Brief sent; questions logged; decision recorded.

Evidence. [BUYER]/Evidence/Team/Executive/CFO/ .

2. Memo

To: [CFO / FD]

From: [Head of AP / Controller]

Re: AP Agent OS — permission, not autonomy

Date: [BUYER]

What this is

We are adopting Evidence Room's AP Agent OS: named agents with charters, human accountability, tests, shadow, and a narrow pilot path. **Proof before permission. Agents earn responsibility. Evidence decides.**

We are not implementing “autonomous AP.” Payment release, vendor bank-change approval, policy exceptions, and legal disputes stay human at every autonomy level.

Why now (operational, not hype)

[BUYER: 4-6 factual bullets – volume, exception ageing, overtime, control findings. No invented industry stats.]

If you want external context only: Deloitte (8 Oct 2025) reported **63%** of organisations had deployed AI and **21%** reported measurable ROI — a reason to **measure**, not to assume return. Hackett (Nov 2025) reported **software-adopter** touchless and productivity figures (60% touchless; 3.5× productivity at ≥30% touchless). Those are **not** our targets and **not** Evidence Room results.

What I need from you

- 1. Confirm sponsor role and that Hold is an acceptable outcome.
- 2. Acknowledge that L3+ or payment-adjacent expansion will come back to you.
- 3. Confirm Controller ownership of any cash/P&L claim (KPI-FIN-SAV).
- 4. Approve funding for [discovery / first increment] of [rounded amount] — **not** a stack-wide transform.
- 5. Confirm we will not use personal or unapproved model tools on invoice data.

How we will talk about money

- Conservative / Base / Upside model — Conservative on the cover.
- Hours released stay **ESTIMATE** until a time study exists.
- Hours become cash only with a mechanism (overtime, contractor, vacancy).
- Duplicate *flags* are not savings. Discounts *available* are not savings.
- False precision is a defect.

[Attach business case ID or “model not yet built – funding is for discovery.”]

First increment (allow-list)

- Entity: [BUYER]
- Channel / type: [BUYER]
- Agent: [e.g. Intake + Duplicate @ L0]
- Verbs: none until Shadow earns them
- Out: payment release; bank write; [add]

How we will know it works

Professional path: labelled historical test → UAT → Shadow (no action) → optional controlled pilot → Measure in four families (risk-control can veto). Promotion is one increment at a time.

Risks I want you to see

Risk	What we will do
Wrong/duplicate payment	Agent 10; human release; no retry-on-timeout writes

Risk	What we will do
Bank details from a PDF	EX-PAY-001; no master-data write tool
Silent model/prompt change	Hash pin; change class C2/C3
Data in unsanctioned tools	Approved list; Privacy review
Date-driven go-live	Gate table; I will bring you a Hold

What this brief is not

A compliance opinion. A fraud-detection claim. A promise that CPI will fall in year one (CPI falls when **cash** cost falls, not when people are “freed” on a slide).

Decision requested

- ☐ Fund discovery
- ☐ Fund first increment as scoped
- ☐ Hold
- ☐ Stop

Signature / date: __

3. What can go wrong

Appendix savings “for colour.” Assistant adds a vendor ROI. Fix: Controller line-out.

4. ACME Manufacturing — ILLUSTRATIVE opener only

Fictional ACME used 8,400 invoices/month, 9 FTE, and a Conservative case that was **net negative** before a labour mechanism — and still funded discovery for control evidence. Do not copy the numbers.

Proof before permission.

End of 02_CFO_BRIEF.md

Evidence Room — AP Agent OS Team Edition

Change management — 01 Toolkit

Product: Evidence Room AP Agent OS — Team Edition

Module: Change Management

Standard: Proof before permission

Audience: Change lead, AP Manager, Head of AP, Communications

Version: 1.0

Change work makes the **operating rules usable** on the floor. It does not “drive adoption” of execution rights the agent has not earned.

1. What to do

Map stakeholders, impacts, resistance, and messages. Time communications to **gates**, not to a launch party. Measure understanding and behaviour (do-not-wait in Shadow), not sentiment about AI.

How. Instruments below.

Who. Change lead Responsible for the plan; Head of AP Accountable for what is said about AP work; AP Manager for floor rhythm.

Wrong. Campaign: “AI frees you for higher-value work” as a promise. Adoption KPIs that reward rubber-stamping.

Control. Language list §6. Comms signed by Head of AP.

Measure. Training before Shadow; incident from “I thought it could pay”; unofficial tool use.

Evidence. [BUYER]/Evidence/Team/Change/ .

2. How — impact map (complete per increment)

Group	What changes in their day	What does change	Risk if we mess up the message	Owner of message
Processors	Extra pane; new codes; review minutes	They still own the live item in Shadow	They wait or obey	AP Manager
Exception owners	Packet quality expectations	They still dispose their family	They dump UNMAPPED	Exception lead
Buyers / GR	Possible template chases later	They still receipt / own PO	They ignore mail as spam	Procurement / GR
Payment authoriser	Maybe a better proposal file	They still release	They think the agent paid	Treasury
Master data	More EX-PAY-001	Dual control remains	They take PDF IBANs	MDM
Controllers	Scorecards	Still sign cash	They see ROI theatre	Controller
Sponsor	Holds are normal	Not a demo factory	Date pressure	CFO + HoAP

3. How — resistance (treat as information)

Pattern	What it often means	Do	Do not
"It will take my job"	No mechanism conversation	Tell the truth: no FTE promise in this OS	Promise safety you cannot keep or promise "no jobs ever"
"We tried OCR"	Prior evidence ignored	Show what is different: gates, not a new engine	Dismiss
"Just go live"	Sponsor pressure	Point at Test pack	Soft launch
Silent workaround	Packet unusable	Fix packet (UAT defect)	More posters
Rubber-stamp AGR	Gaming autonomy	Sample QA; intervention KPI	Celebrate low intervention

4. How — stakeholder cadence

Audience	Before workshop	Before Shadow	Before Pilot verbs	After Measure
Floor	Questionnaire	Training + M6 briefing	Verb list + examples	What was evidenced
Control/IA	Interview	Access to folders	Rollback	Pack
Sponsor	CFO brief	"Shadow ≠ live"	Allow-list	Conservative results
Adjacent (buyer/GR)	If in slice	Only if they will be chased	Template + who to call	Stop/continue

5. How — manager conversation (15 min script)

AP Manager to team:

1. Why this slice (operational fact).
2. What the agent **cannot** do this month.
3. How you will be measured (not "AI usage").
4. How to reject.
5. How to raise an incident.
6. "If I ever tell you to follow the agent instead of the SOP, that is a defect — tell Control."

6. Language list

Use

- Proof before permission.
- Agents earn responsibility.
- Evidence decides.
- Shadow is not an instruction.
- Hold is a valid outcome.
- ESTIMATE (for hours).
- Human release.

Do not use

- Autonomous AP / autopilot / set-and-forget.
- Guaranteed savings / ROI / "2× productivity."
- Fraud-proof / fraud detection as a product claim.
- Compliance guaranteed / audit-ready as a certificate.
- World-class cost (unless a dated buyer-owned study).
- "The AI decided."
- Industry statistics you did not source (and never invent).

External research, if needed: label Deloitte 8 Oct 2025 (63% deployed / 21% measurable ROI) and Hackett Nov 2025 **software-adopter** touchless figures as **not our results**.

7. What can go wrong

Failure	Control
Change success = adoption %	Ban; use override quality + FNR
Comms outrun the charter	Head of AP sign-off
Floor hears a different story than Steer	Same facts, different length

8. Measure change work

Question	Metric
Did people hear the rules?	Training pass before Shadow
Did they wait?	Shadow contamination notes
Did they invent authority?	Incidents tagged messaging
Did we over-claim?	Language review of sent comms

Evidence. Plan, signed messages, review log.

Proof before permission.

End of 01_CHANGE_TOOLKIT.md