

90-Day Organic Content System

Product	AP Agent OS
Version	1.0
Date	September 2026
Use	Implementation toolkit

Responsibility is earned.

Mix

80% useful. 20% commercial. Primary: LinkedIn. Secondary: newsletter, SEO pages, occasional Shorts/X.

Pillars: AP operations · Finance AI · Agent governance · Transformation economics · Controls · AP metrics · Implementation lessons · Process redesign.

30 LinkedIn posts (titles only; first five fully written)

1. Approvals take too long for 49% of AP leaders (Arden 2025). That is not an LLM problem first.
2. Exception rate 18.4% average. Code them or keep folklore.
3. STP 35.4% is not a moral failing. It is a design gap.
4. Best-in-Class \$2.65 vs all others \$12.42 per invoice. Know which table you sit in before you buy a copilot.
5. 44% of AP teams use some AI. 22% have no plan. The missing artefact is a charter, not a licence.
6. Invoice text is data, not instructions.
7. If your agent can release payment, you do not have an agent. You have an incident.
8. Deterministic match engines still beat prose at three-way match.
9. GRNI is an ownership problem dressed as a systems problem.
10. Near-duplicate is a review queue, not a fraud headline.
11. DOA in a PDF is not machine-readable control.
12. A weekly "AI invoices processed" tile is a vanity metric.
13. Shadow mode is a control, not a delay tactic.
14. Promote agents the way you promote people: evidence, then scope.
15. Supplier emails drafted by models still need a human name.
16. Bank-change + payment proposal is the red team scenario.
17. Shared services without a taxonomy cannot compare sites.
18. The first RACI mistake is putting the orchestrator in the Accountable cell.
19. Cost per correct outcome beats cost per token.
20. If Internal Audit first hears about the agent in a demo, you are late.
21. Paper is not romantic. It is an intake channel with a risk profile.
22. A 4–6 week agent is a bounded story, not a promise.
23. Override rate is a control metric. Celebrate a good reject.
24. Root cause without a test is a story.
25. Month-end completeness is a list with IDs, not a narrative.
26. Copilot in the ERP is a tool. Who writes the job description?
27. You cannot govern sixteen agents if you cannot name one owner.
28. Lost early-pay discounts are only a benefit if you measure them now.
29. The kill switch must work on a Wednesday when the owner is off.
30. Responsibility is earned. (Soft CTA to diagnostic.)

Post 1 full: Arden Partners' 2025 ePayables study (204 AP organisations) puts "approvals take too long" at 49% of top challenges. Before you add an approval copilot, publish the stall report: whose queue, how old, what value, whether a delegate exists. Agents are useful here at Level 1. They should not approve.

Post 2 full: 18.4% exception rate is the average in that same study. If your codes are free text, you cannot run a root-cause agent. Start with ten codes. Missing PO, price, qty, missing GR, duplicate, tax, approval, coding, aged, system.

Post 3 full: Straight-through processing at 35.4% average vs 51% Best-in-Class. The gap is usually PO discipline, supplier enablement and GR — not a missing prompt. Map those three before you staff an orchestrator.

Post 4 full: \$9.84 average all-in cost per invoice. \$2.65 Best-in-Class. \$12.42 everyone else. Use these as orientation. Your number needs your labour and your volume. I built a conservative model that is allowed to say “don’t spend yet”.

Post 5 full: 44% of AP teams already use some form of AI. The failure mode I see is uncharted use: someone pastes invoices into a consumer tool. Write the approved-model list first. Then write the agent.

10 carousels

1. Five autonomy levels
2. Sixteen agents on one slide
3. Exception taxonomy (10 codes)
4. Human / deterministic / agent
5. Control matrix excerpt
6. KPI four families
7. 10-step method
8. Payment proposal red lines
9. Monday sequence
10. Prompt-injection on an invoice

5 long-form articles (SEO)

1. AP agents vs AP automation — the missing operating layer
2. How to write an AP agent charter
3. Exception taxonomy for shared services
4. Measuring agents without vanity metrics
5. A control framework for Finance agents (NIST/COSO mapped in plain English)

10 executive charts

Cost/cycle/exception/STP (cite Ardent); BIC vs others; AI adoption 44/34/22; challenge bars; first-agent quadrant; autonomy ladder; heatmap; payback sensitivity; queue ageing; promotion gate.

10 diagnostic posts

One per domain question cluster (volume, exceptions, process, data, controls, measurement) plus four “score yourself 0–3” posts.

10 practical frameworks

Observe-extract-structure; agentisation matrix; RACI with agents; gold-set design; shadow rules; SoD for operators; hostile-content handling; benefits rules; 90-day plan; kill switch drill.

Newsletter

Monthly: one benchmark, one artefact, one caution. No launch-week spam after day 14.

This material is an implementation toolkit. It is not legal, tax, accounting, audit or investment advice. Agents do not authorise payments. Humans remain accountable. Statistics are cited in the Research Ledger.