

Evidence Room

Agents that earn responsibility.

AP AI Readiness — Free Diagnostic Guide

Score your Accounts Payable operating posture before you grant agents any authority.

SKU ER-FREE-AP · Educational kit · 8–12 pages

v1.0 · Commercial pack · Not legal advice · Not a software runtime

1. Purpose

Most AP AI efforts fail because teams skip the operating model and jump to prompts. This guide exists so Controllers and AP leaders can score readiness honestly, decide what agents may do, and refuse unsafe autonomy.

It accompanies the Free AP AI Readiness Kit: diagnostic, maturity model, opportunity heatmap, ten-agent overview, KPI worksheet, business-case starter, and scorecard. The output is a posture — not a vanity score.

2. How to run the diagnostic

- Schedule 60–90 minutes with an AP lead present; invite Controls and IT for access, SoD, and integration items.
- Score each of the 35 statements 0–4 using the rubric. Prefer evidence over optimism.
- Require an evidence note for any score of 3 or 4. If you cannot point to a control, report, or SOP, lower the score.
- Sum each dimension, then total out of 140. Map to maturity levels before discussing any pilot.
- Agree two baseline KPIs and one control KPI that must not worsen before any agent shadow mode.

3. Scoring rubric (0–4)

Score	Meaning	Evidence expectation
0	Absent / unknown	No artefact, no owner
1	Ad hoc	Exists in pockets; tribal knowledge
2	Emerging	Documented but unevenly followed
3	Established	Mostly followed; measurable gaps
4	Managed	Consistent, reviewed, improvable with evidence

4. Dimension deep-dive — Process & Data

Process (7 questions · max 28)

Intake channels and owners, completeness standards before queue entry, separation of PO and non-PO paths, exception catalogues with owners, current approval matrices, payment calendars with cut-offs, and checklist-driven close tasks. Low process scores mean agents will accelerate confusion.

- Document every intake channel before automating capture.
- Catalogue exception types before building triage agents.
- Do not pilot matching if GR/PO hygiene is unknown.

Data (6 · max 24)

Vendor master payment fields, duplicate vendor reviews, clean ERP field mapping, joinability of PO/GR/invoice IDs, ≥12 months historical queryability, and logged remediation. Agents inherit your data defects; they do not cure them.

5. Dimension deep-dive — Controls, Technology, People, Measurement

Controls (7 · max 28)

SoD between vendor setup, entry, and payment release; dual control on bank changes; audit trails; limited force-pay; annual access reviews; monitoring of high-risk classes; and written discount/hold policies. Controls gate earn-up — they are not optional polish after the demo.

Technology (6 · max 24)

System of record (not spreadsheet-as-ledger), capture/OCR share, workflow beyond email chase, stable procurement/receiving integrations, non-heroic reporting extracts, and a sandbox for pilots.

People (5 · max 20)

Clear roles, ability to explain approvals from system evidence, redesign capacity (not only firefighting), leadership stance that agents sit under human accountability, and measured training adoption.

Measurement (4 · max 16)

Cost per invoice (or proxy), exception and STP rates with agreed definitions, cycle time by class, and fixed-cadence KPI reviews with owners and actions.

6. Maturity levels → agent posture

Level	Points	Label	Safe agent posture
1	0–27	Manual & opaque	Observation only — map processes and KPIs
2	28–55	Documented but fragile	Draft-only; 100% human action; no write-backs
3	56–83	Controlled operations	Propose-and-confirm; mandatory human for money/master data
4	84–111	Measurable & integrable	Bounded autonomy for narrow low-risk classes with sampling
5	112–140	Agent-ready with governance	Portfolio with RACI, evidence standards, continuous monitoring

7. Earn-up stages

Responsibility is earned, not asserted. Promote agents only through evidence gates:

- **Observe** — read-only signals; no outbound action.
- **Recommend** — ranked suggestions; human decides.
- **Draft** — artefacts prepared; human edits and sends/posts.
- **Execute (narrow)** — bounded actions on pre-approved classes after earn-in.
- **Expand** — widen scope only with sampled performance and kill-switches intact.

Demotion is a first-class control. Precision slips, SoD near-misses, or evidence gaps return an agent to Draft or Observe.

8. Baseline KPI worksheet

Agree definitions in writing. Capture last full quarter. Industry averages are context only.

KPI	Definition	Your value	Period
Monthly invoice volume	Invoices completed		
Cost per invoice	Loaded AP cost ÷ invoices		
Exception rate	Manual intervention after capture ÷ invoices		
STP / touchless	No human touch after intake ÷ invoices		
Median cycle time	Receipt → payment-ready (days)		
First-pass match	Clean PO match ÷ PO invoices		
Duplicate suspect rate	Flagged potential dups ÷ invoices		
On-time payment	Paid by due/discount date ÷ invoices		

Industry context (Ardent Partners 2025 survey averages): cost/invoice \$9.84 · exception 18.4% · STP 35.4%. External benchmark — not a guarantee of your results.

9. Opportunity heatmap — how to read it

After scoring, plot agents against pain × readiness. Prefer high-pain, high-readiness cells for pilots. High-pain, low-readiness cells need foundation work first. Low-pain cells are distractions, not pilots.

10. What not to do

- Do not skip Level 1–2 foundations to chase automation theatre.
- Do not grant payment release or bank-edit authority to any agent.
- Do not market agents as guaranteed fraud prevention.
- Do not treat illustrative ROI models as forecasts.
- Do not confuse Ardent survey averages with your targets.

11. Recommended product path

Maturity	Recommended next step
1–2	Stay on Free kit; optionally Starter for operating-model literacy
3	AP Agent Starter → selective Pro blueprints
4	AP Agent OS Pro + Custom for edge cases
5	Team change packs + Custom + multi-process roadmap (still AP-first)

12. Session close checklist

- Scorecard completed with evidence notes
- Maturity level agreed (not negotiated upward without evidence)
- Two primary KPIs + one control KPI named
- One candidate agent identified or explicitly deferred
- Internal Audit / Controls informed if a pilot is planned

Disclaimer

Evidence Room provides methodologies, diagnostics, and design artefacts. We do not guarantee cost reduction, ROI, fraud prevention, or audit outcomes. Industry benchmarks attributed to Ardent Partners (2025) — average cost per invoice **\$9.84**, exception rate **18.4%**, STP **35.4%** — are third-party reference points only and are not Evidence Room performance claims. Agents must operate under your controls, policies, and human accountability. Evidence Room is not your ERP vendor, bank, or auditor.

AI agents can be wrong. Humans remain accountable for approvals, postings, payments, supplier master changes, and representations to auditors and regulators.

Appendix — Evidence pack checklist (1)

Every material agent decision should leave an evidence pack an auditor can follow.

This appendix reinforces Evidence Room operating discipline. It is educational and does not modify your licence, create warranties, or guarantee operational outcomes.

- Inputs used and source system references
- Rule or model version hash
- Confidence or variance table
- Human override identity and rationale
- Timestamp and Invoice Case ID
- Kill-switch status at time of action

Remember: Observe → Recommend → Draft → Execute (narrow) → Expand. Demotion is hygiene. Ardent Partners (2025) benchmarks are third-party survey averages for context only.

Operating note 1: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Operating note 2: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Operating note 3: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Operating note 4: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Appendix — Kill-switch protocol (2)

A kill-switch is useless if nobody has rehearsed it.

This appendix reinforces Evidence Room operating discipline. It is educational and does not modify your licence, create warranties, or guarantee operational outcomes.

- Named operators
- Disable write paths first
- Notify AP Manager / Controls / IT
- Record demotion stage
- Preserve logs
- Retest after prompt/model change

Remember: Observe → Recommend → Draft → Execute (narrow) → Expand. Demotion is hygiene. Ardent Partners (2025) benchmarks are third-party survey averages for context only.

Operating note 1: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Operating note 2: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Operating note 3: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.

Operating note 4: maintain named owners for agent, control, and KPI. Refuse unsupervised payment release. Keep evidence packs complete. Review overrides weekly. Treat silence in logging as a defect.